

Basics of AI: what you need to know

Dr. Katie Stolee
Associate Professor
Department of Computer Science
North Carolina State University

Computer Science is Everywhere

Education

Energy and
Sustainability

Scientific
Discovery

Security and
Privacy

Transportation

Technology for
Development

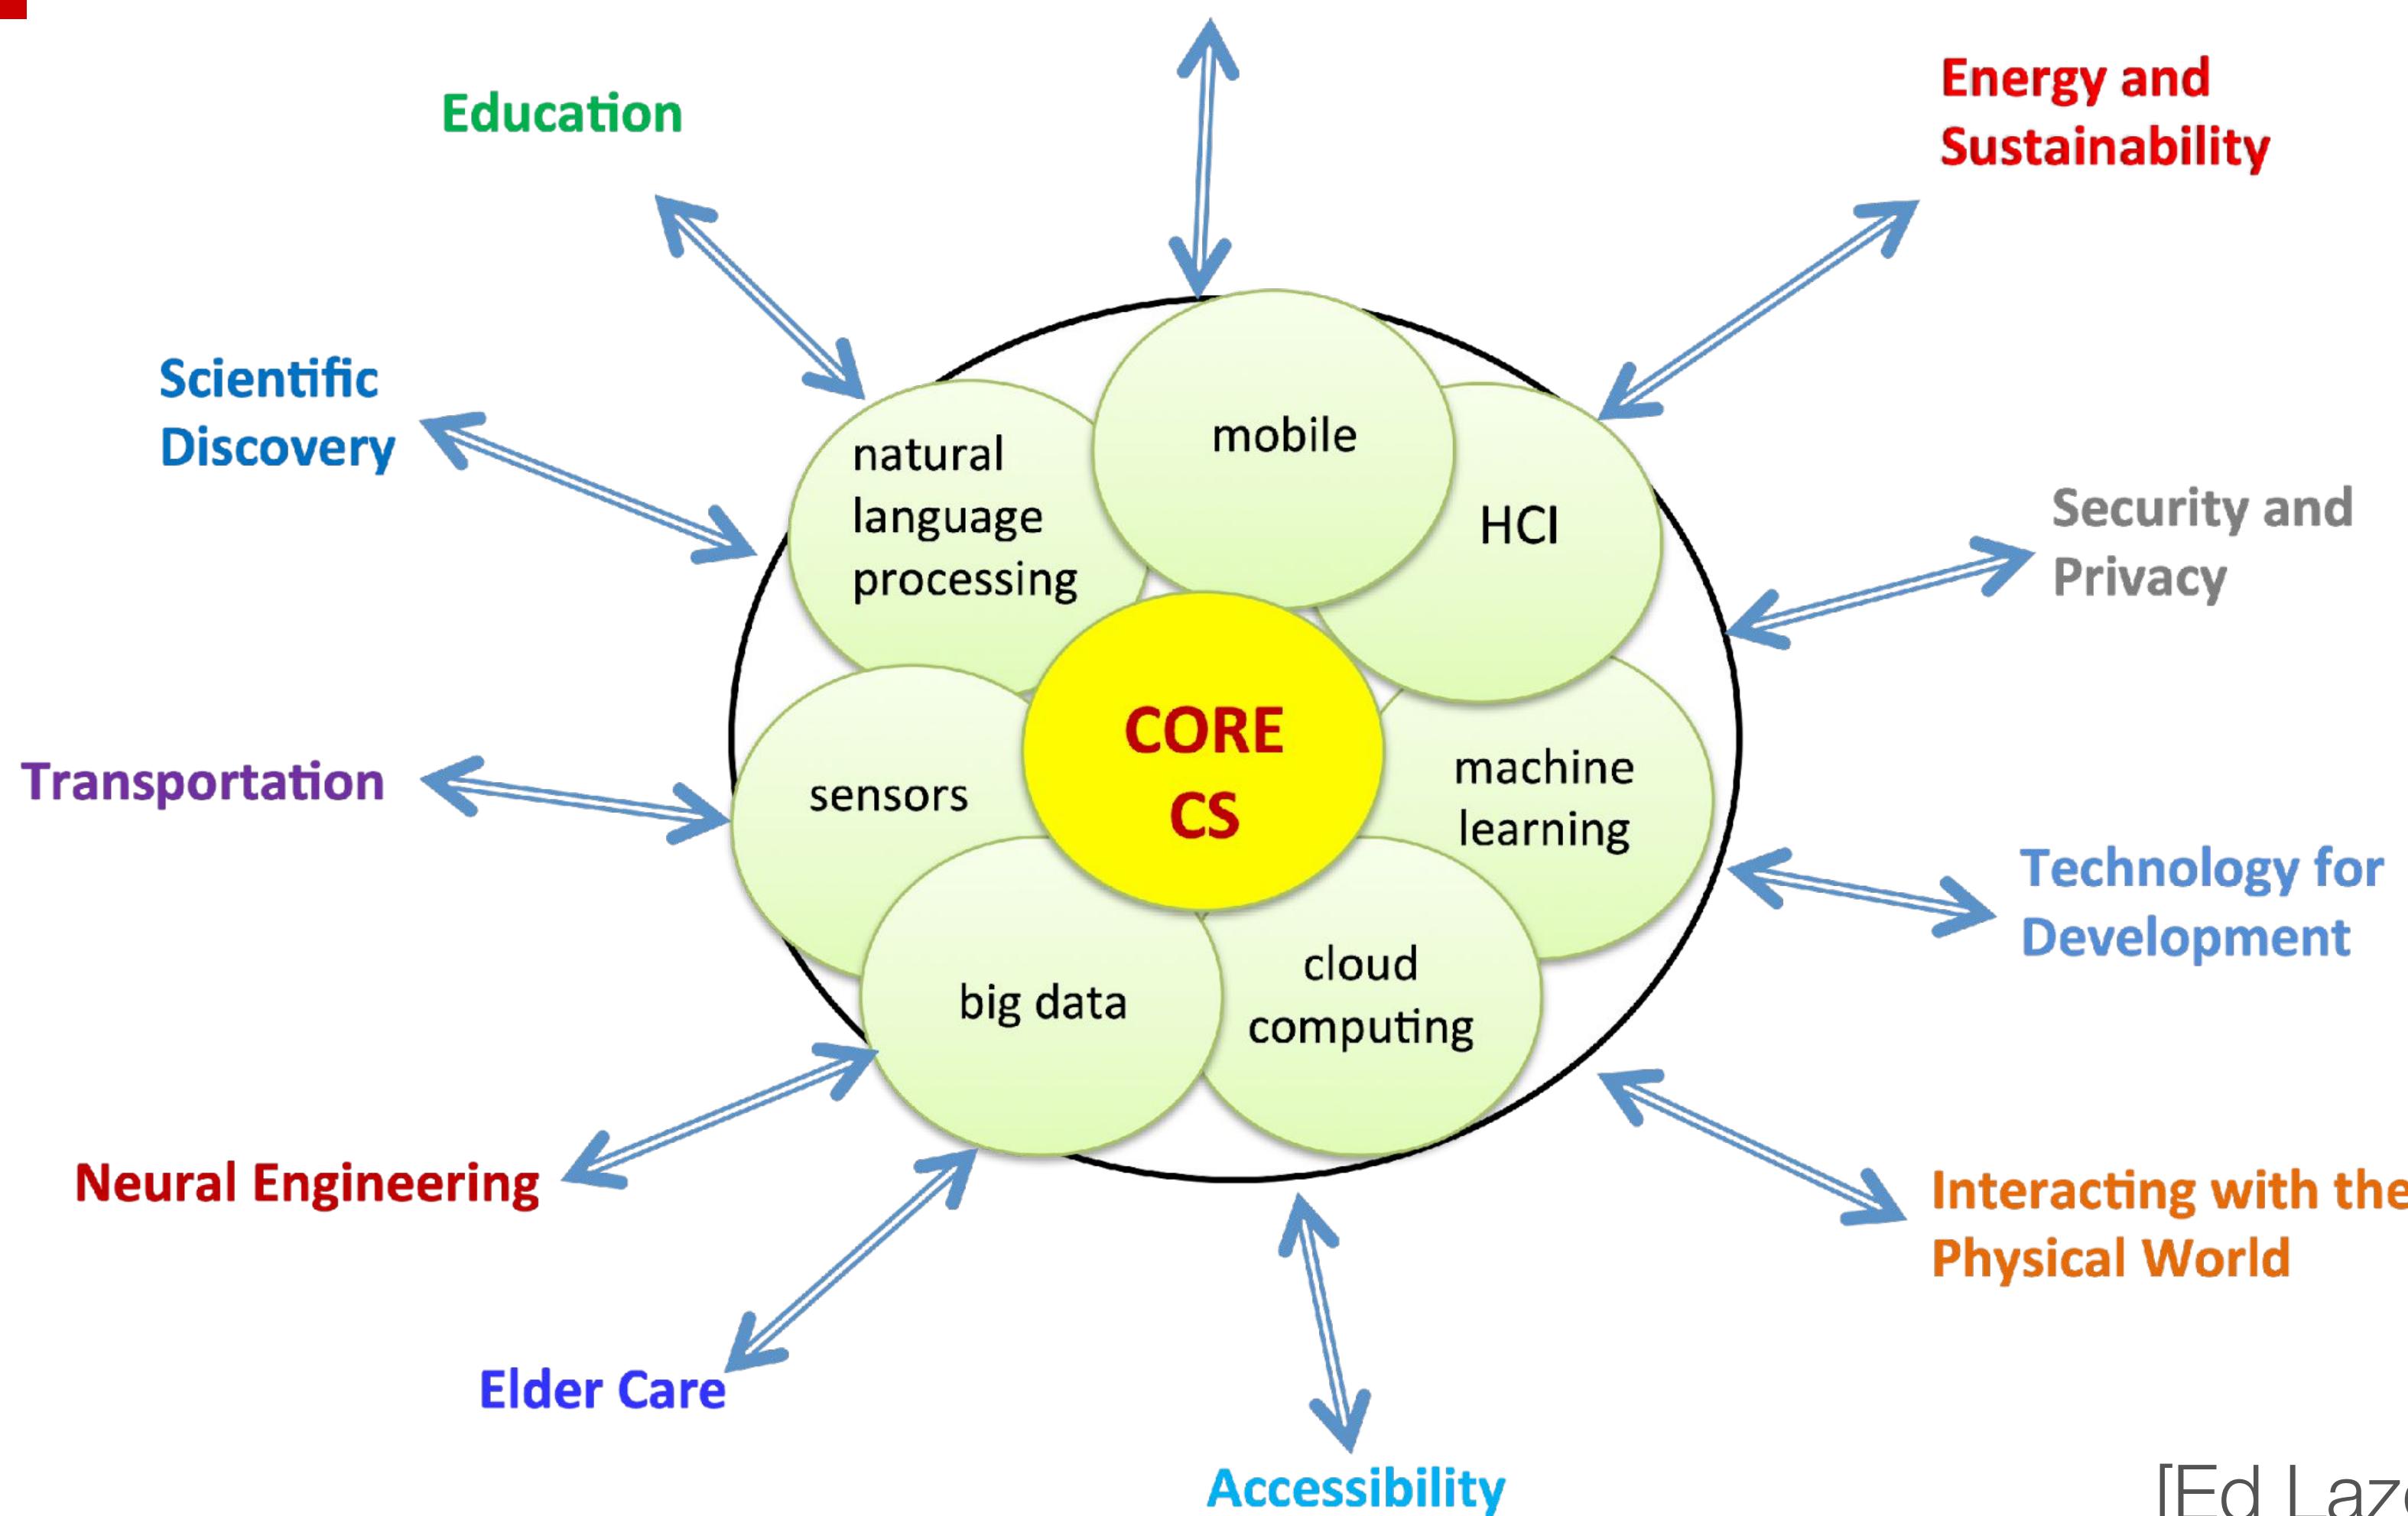
Neural Engineering

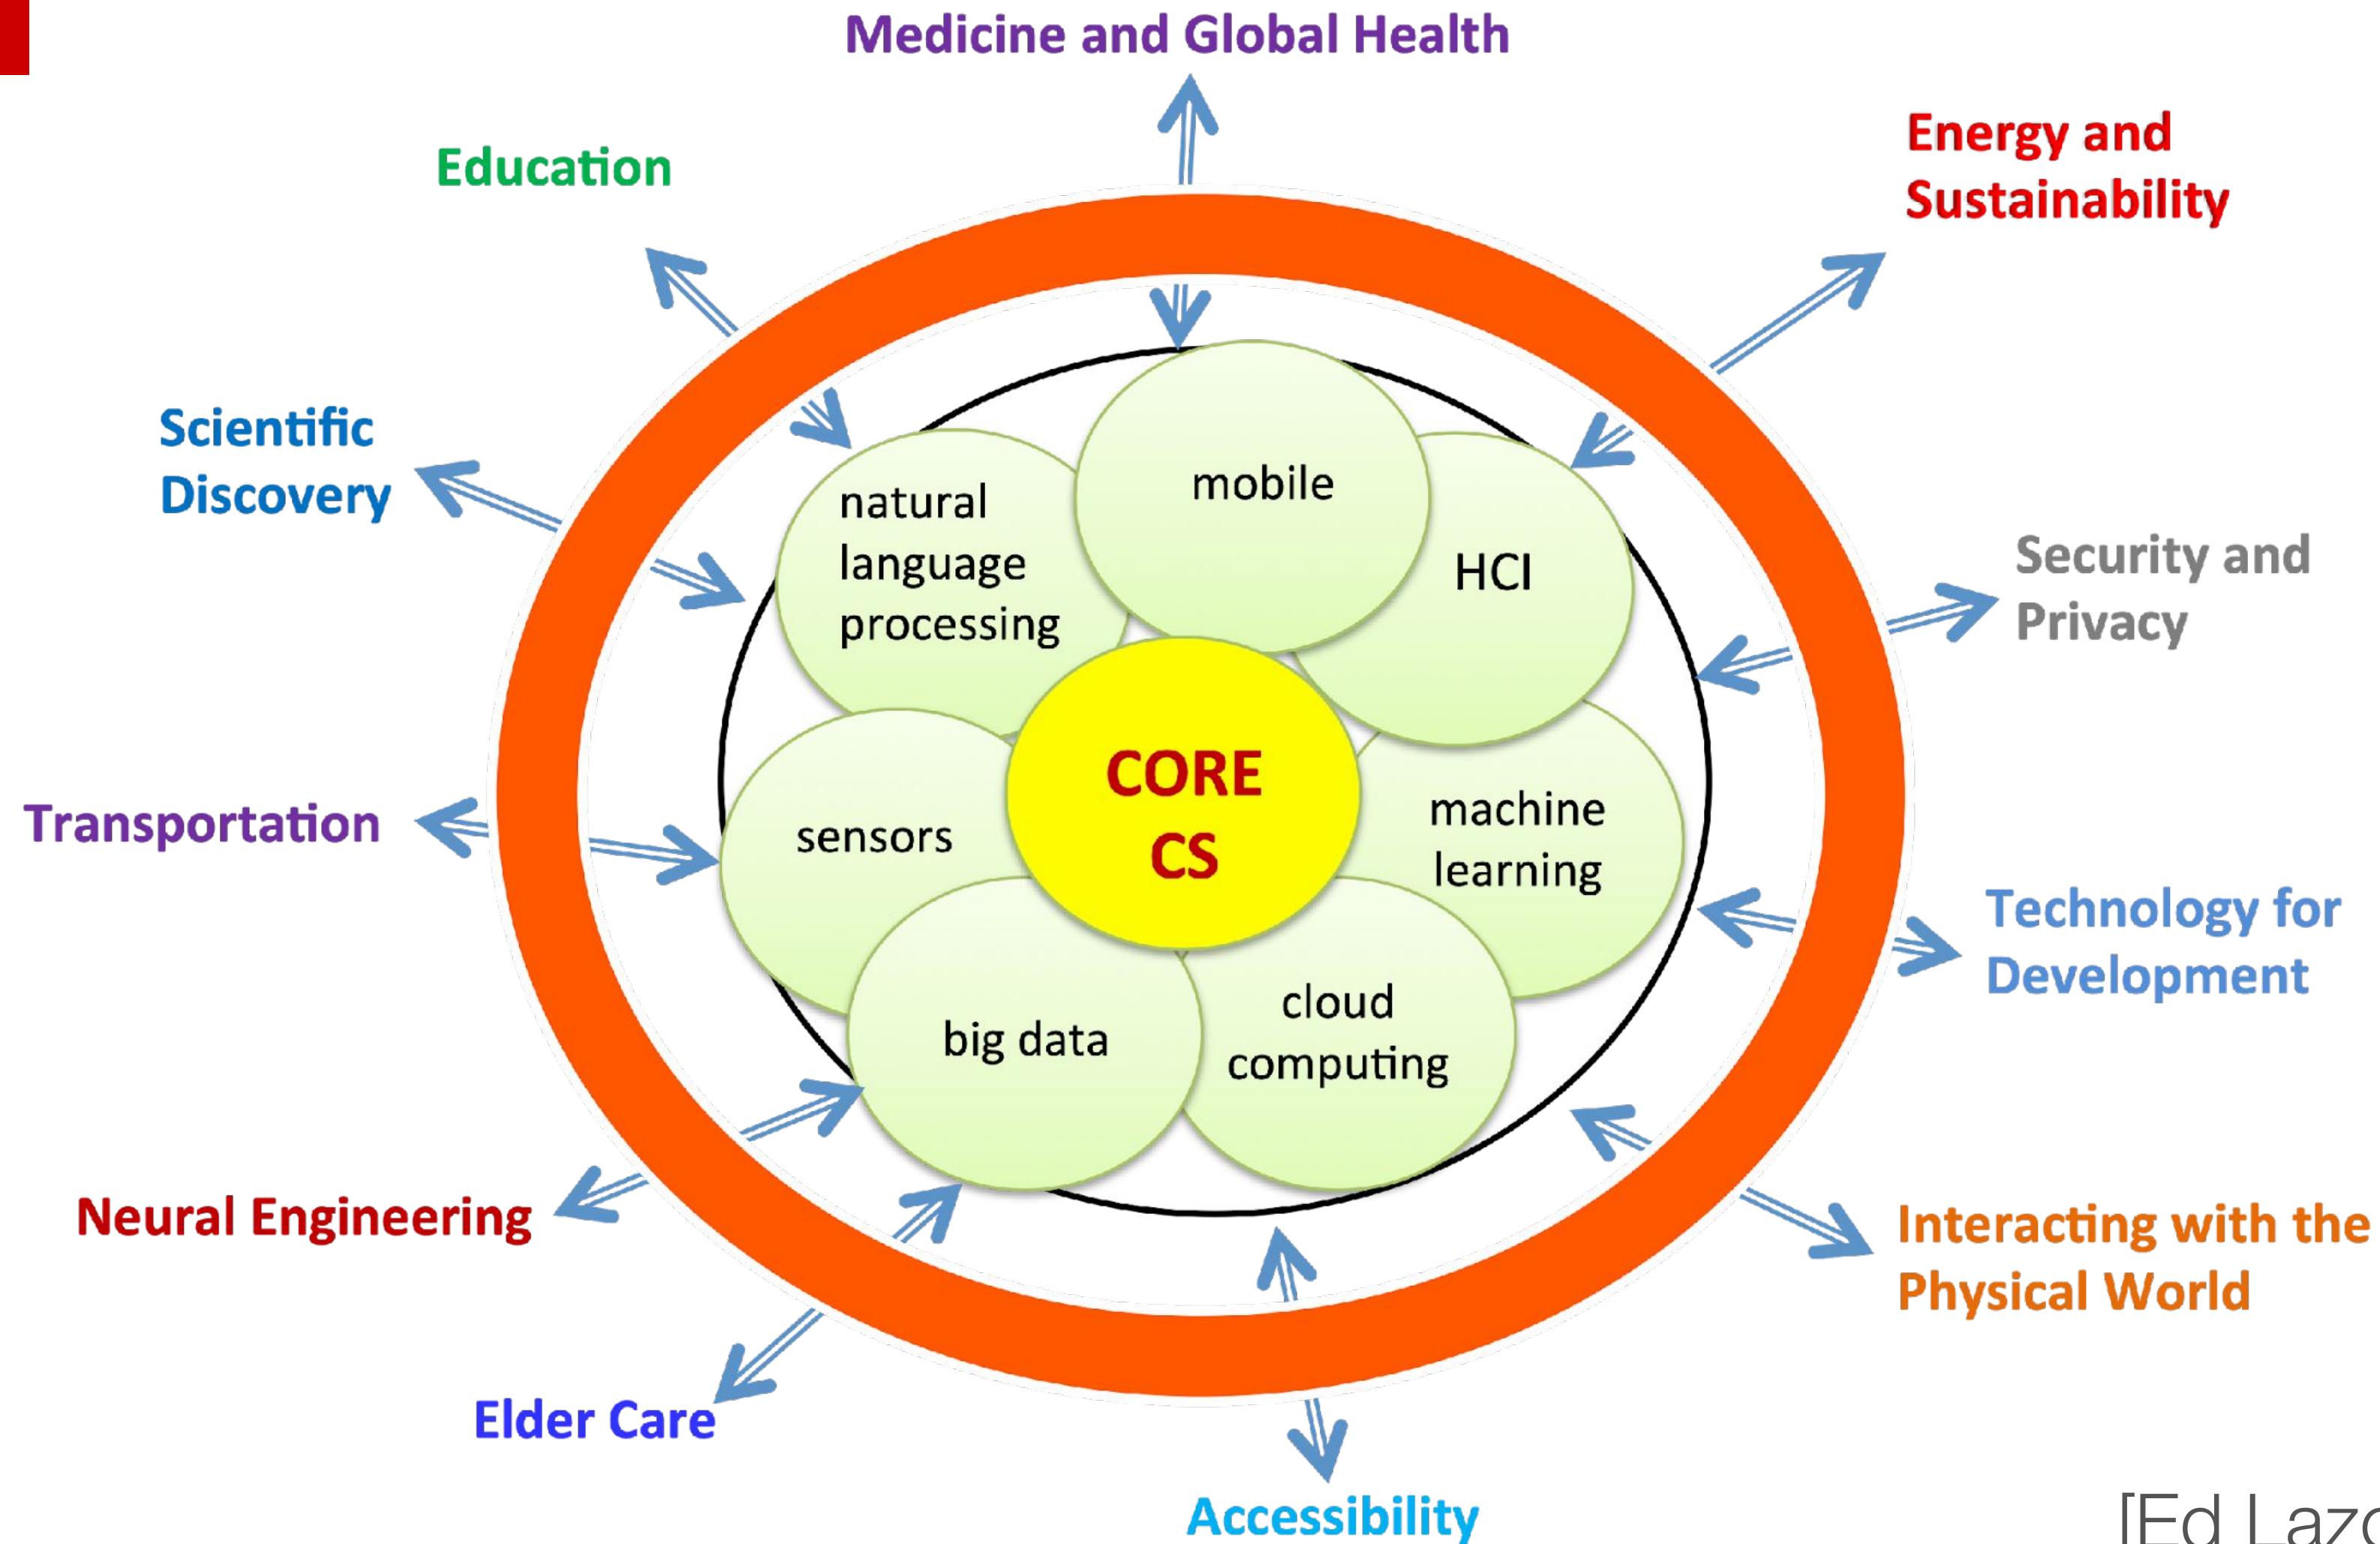
Interacting with the
Physical World

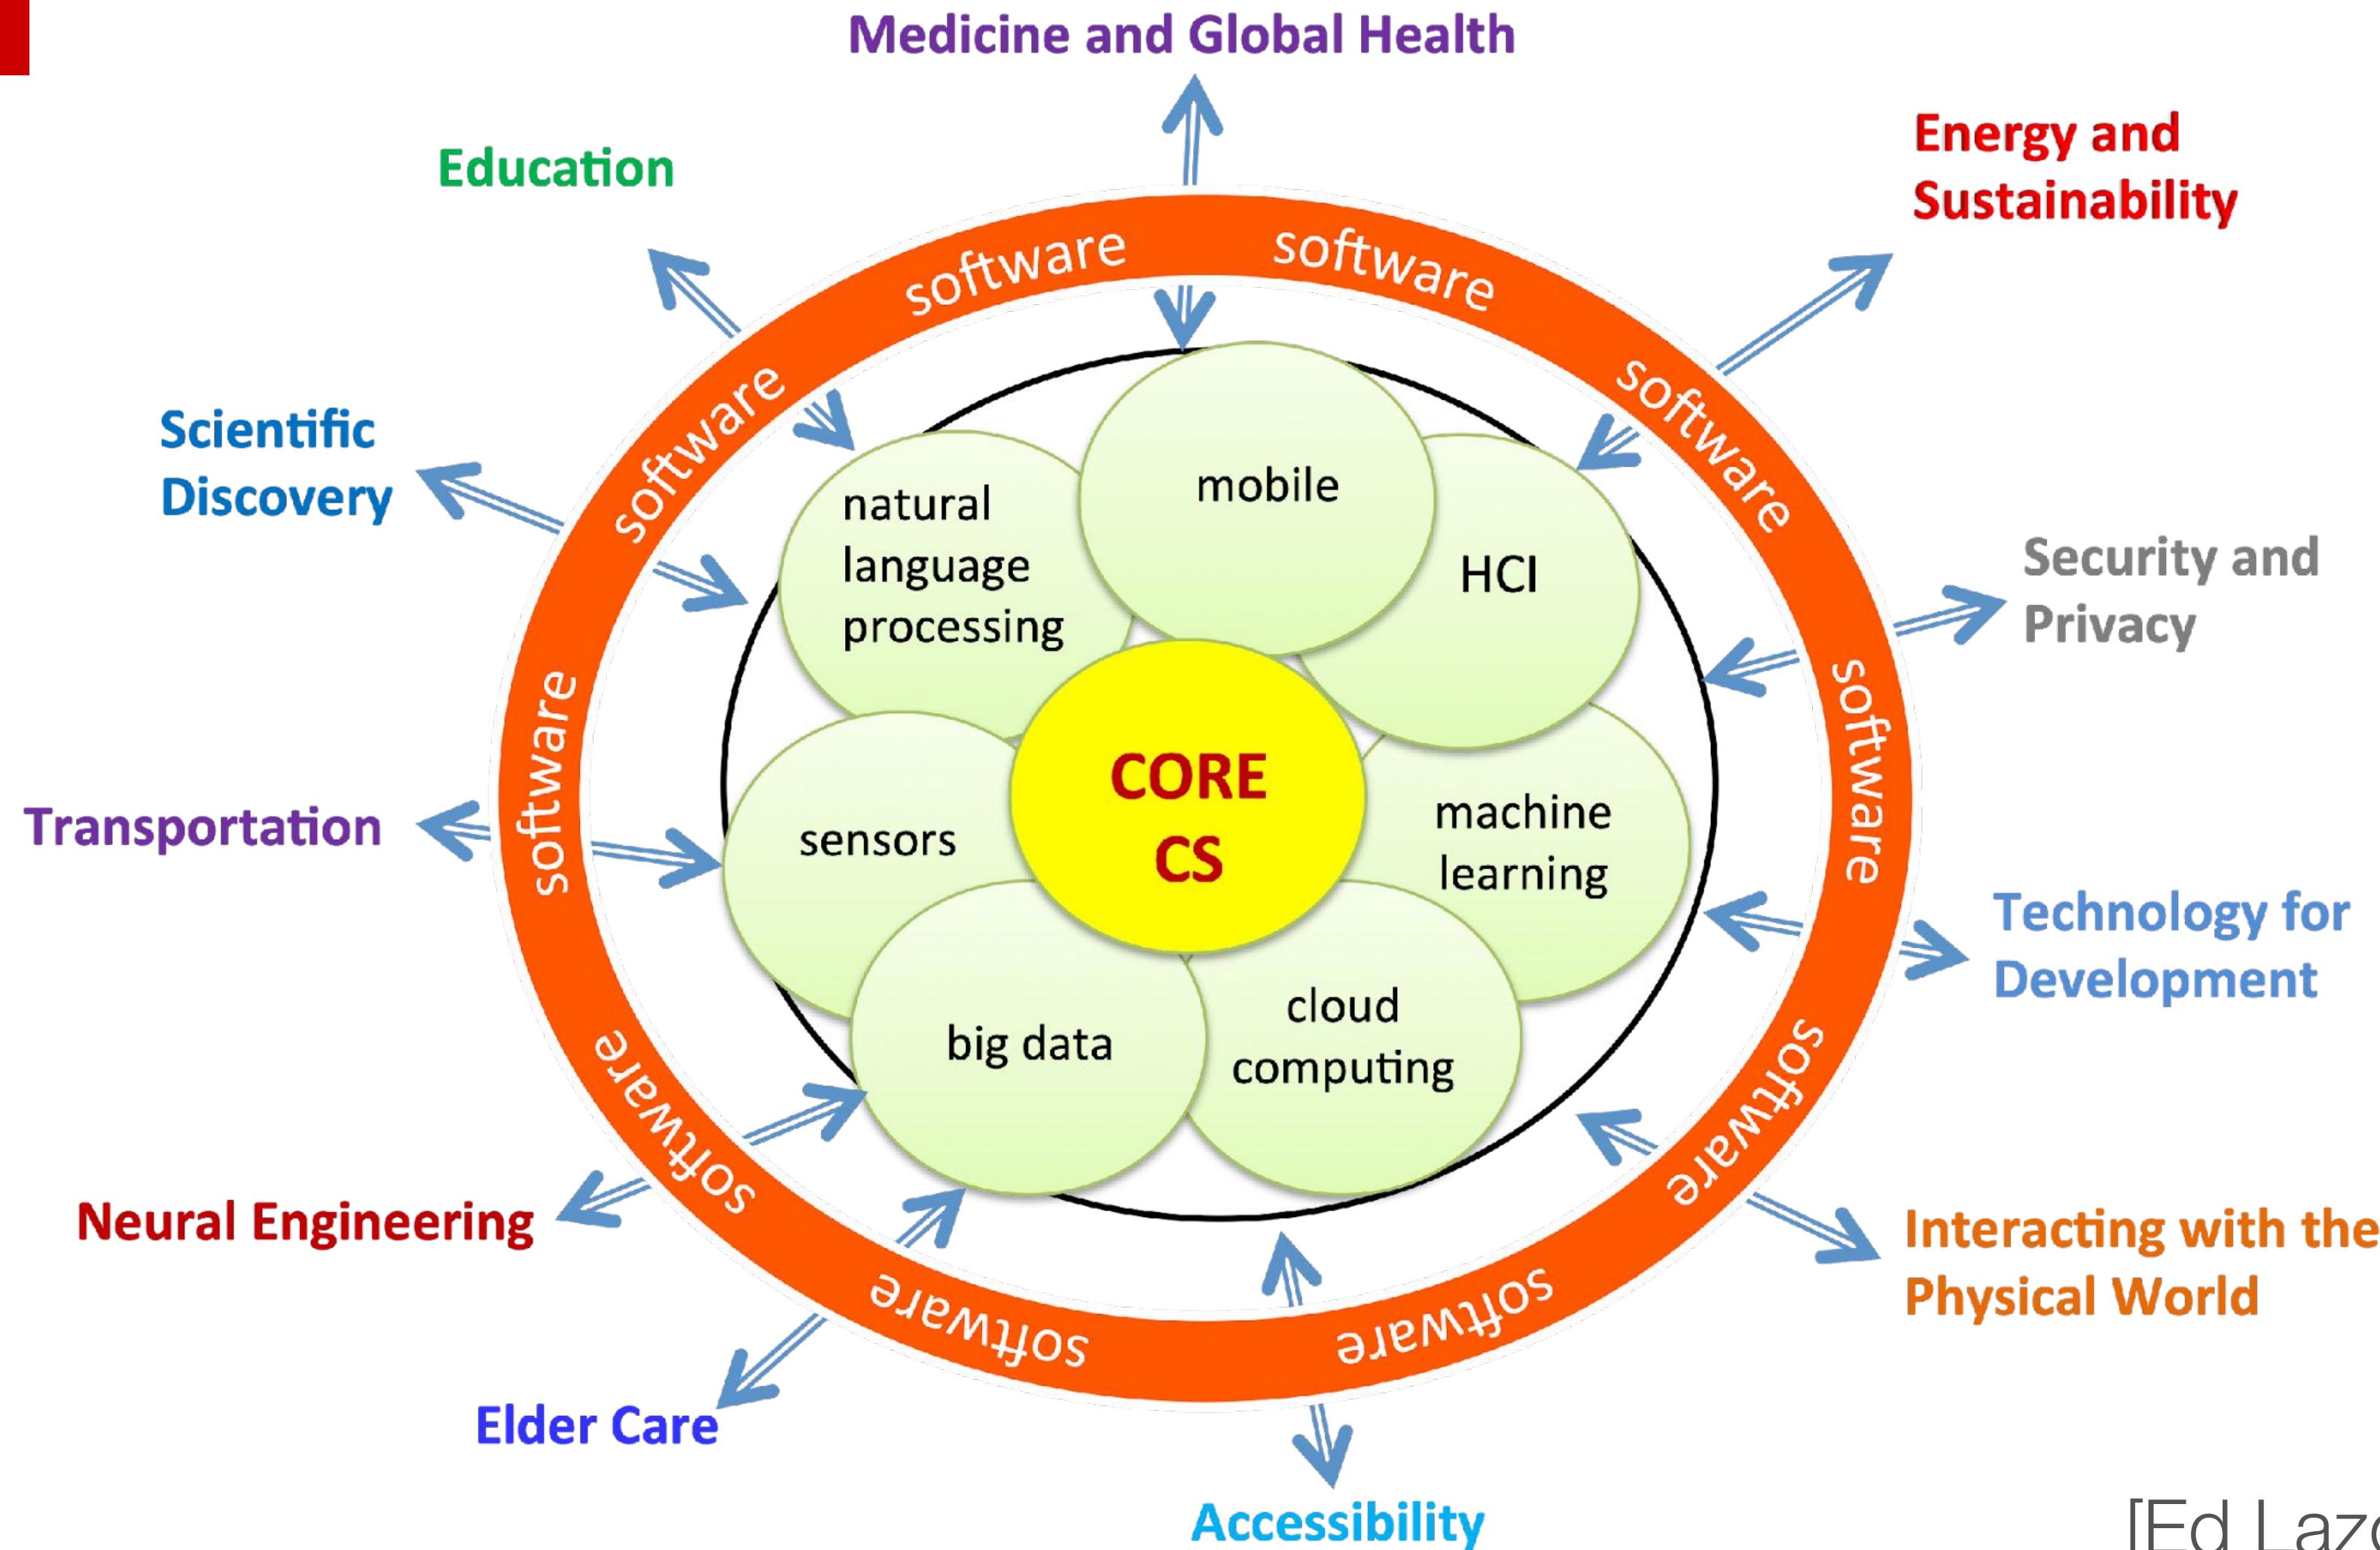
Elder Care

Accessibility

[Ed Lazowska]



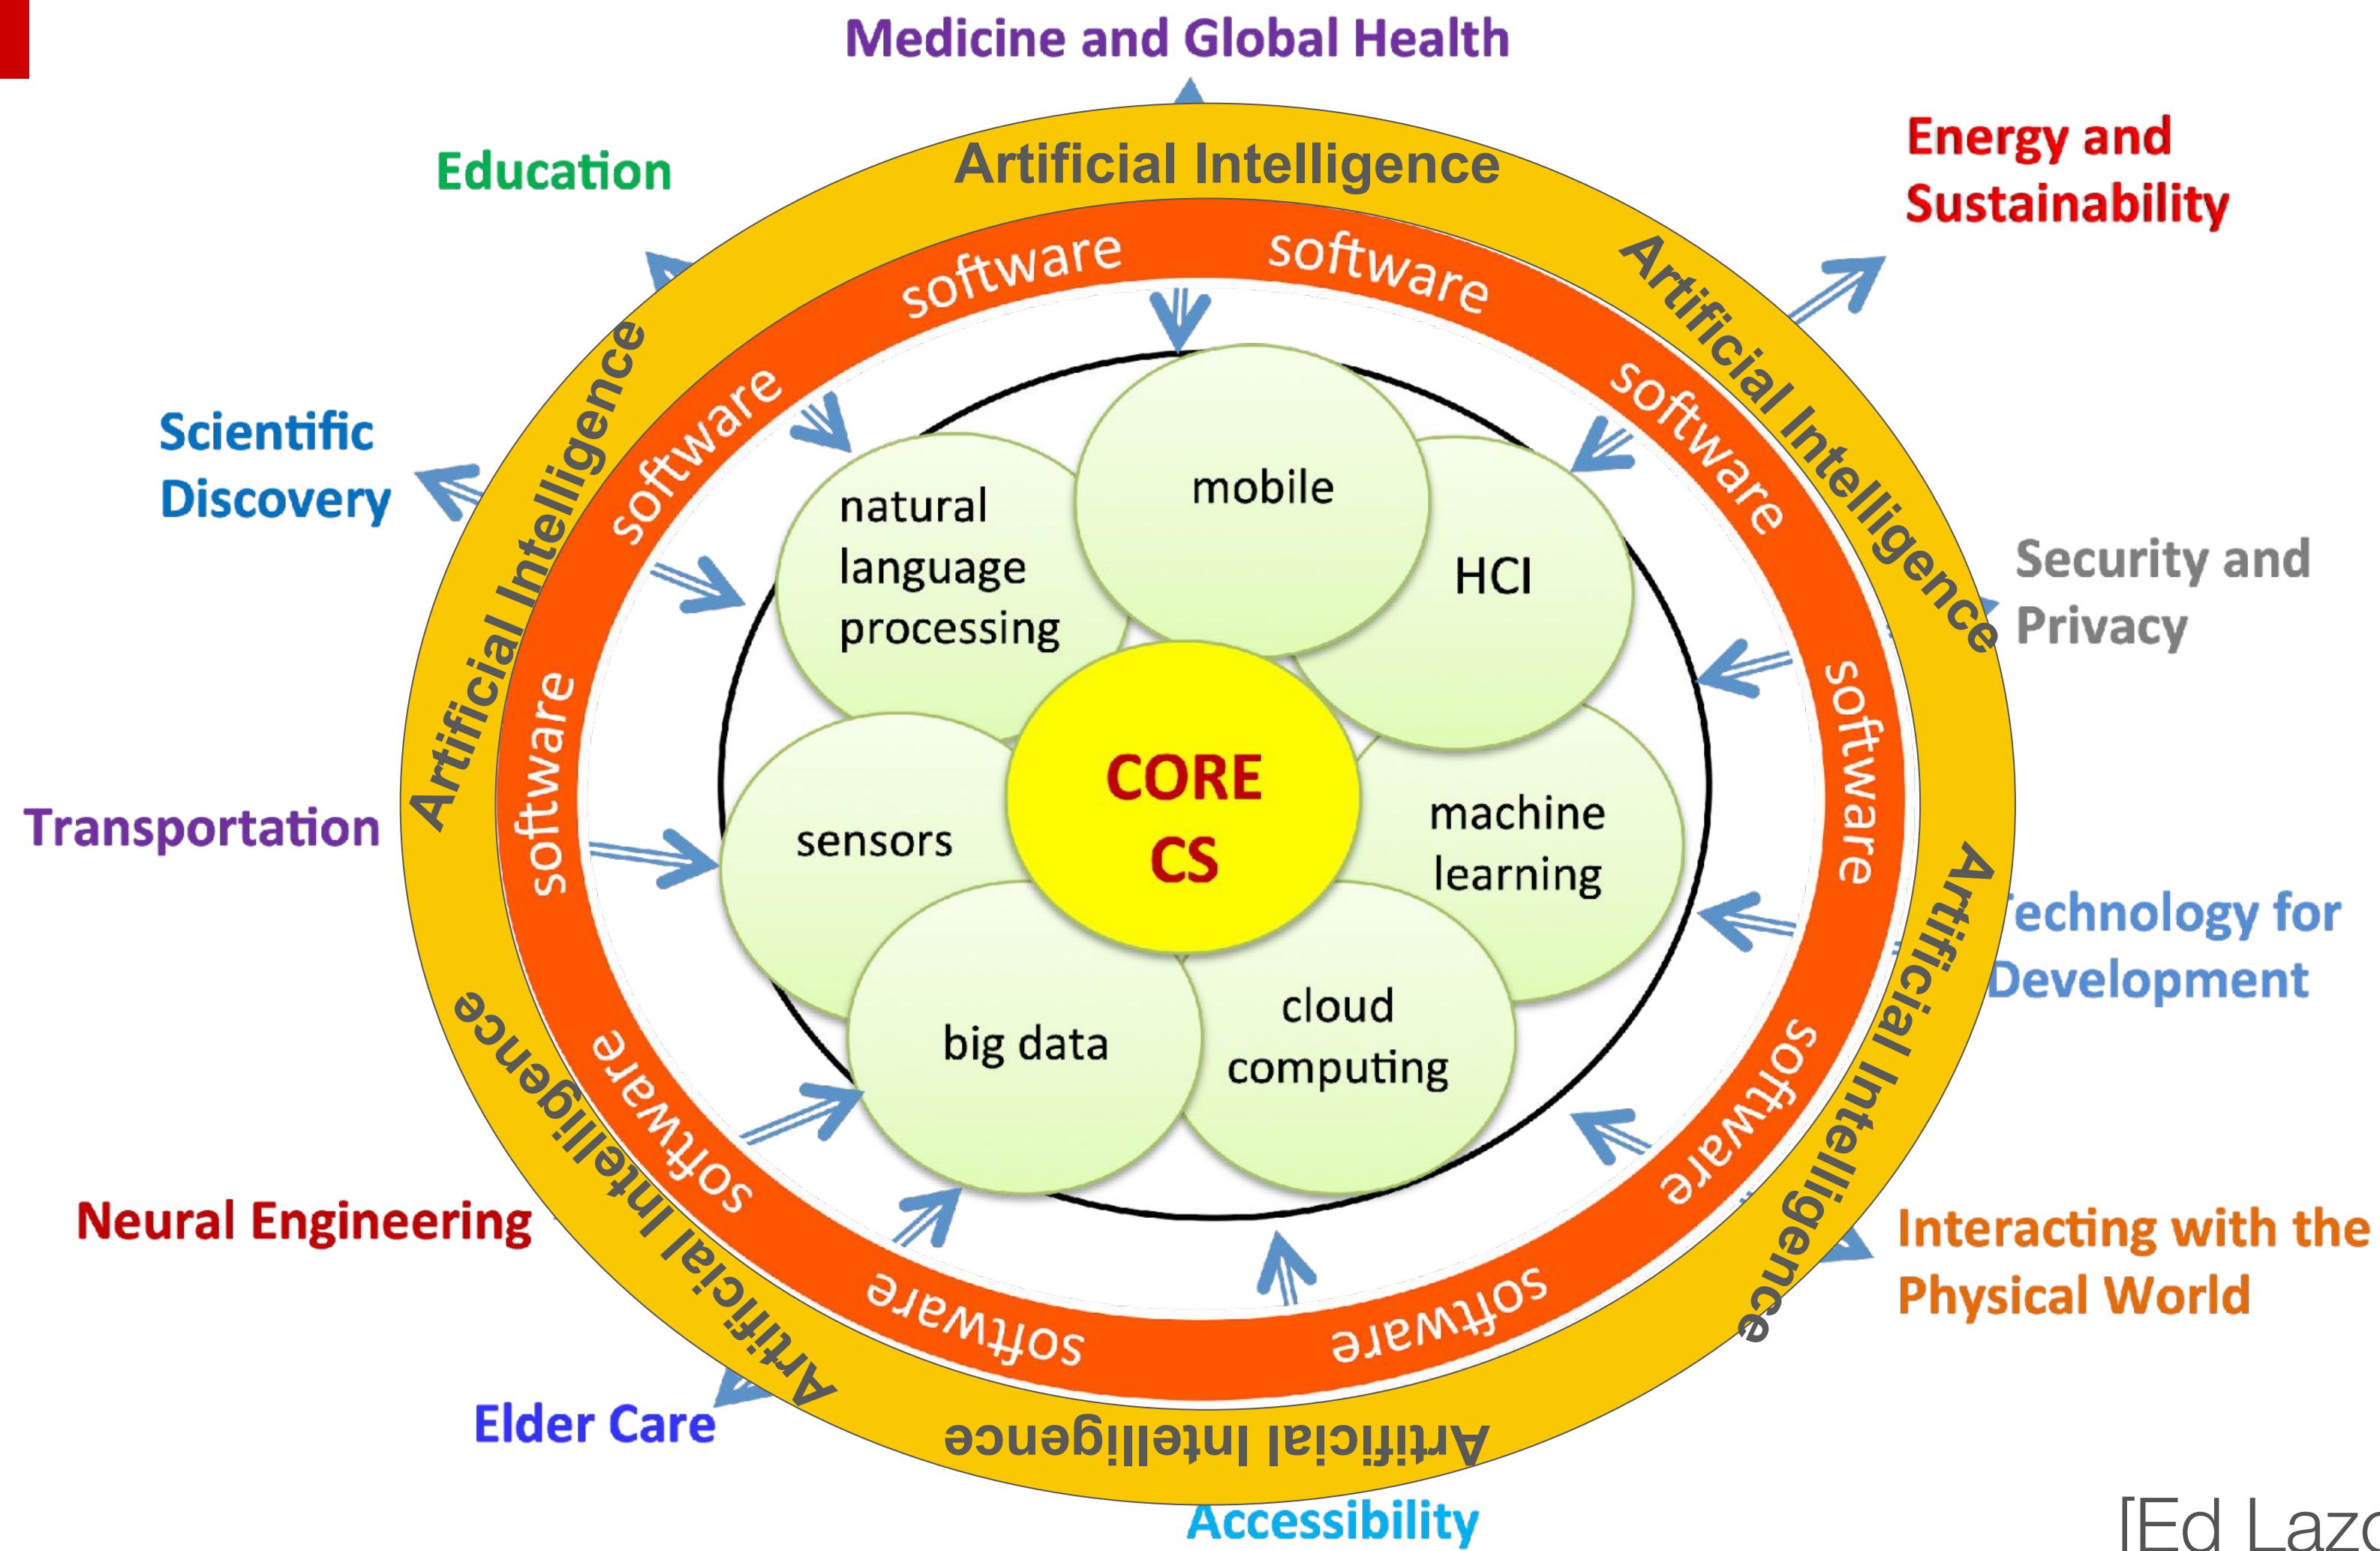




NC STATE

your
life
depends
on
code





What do you know about AI?

Artificial Intelligence

Machines that mimic human cognition

Building machines that can solve
problems and make decisions



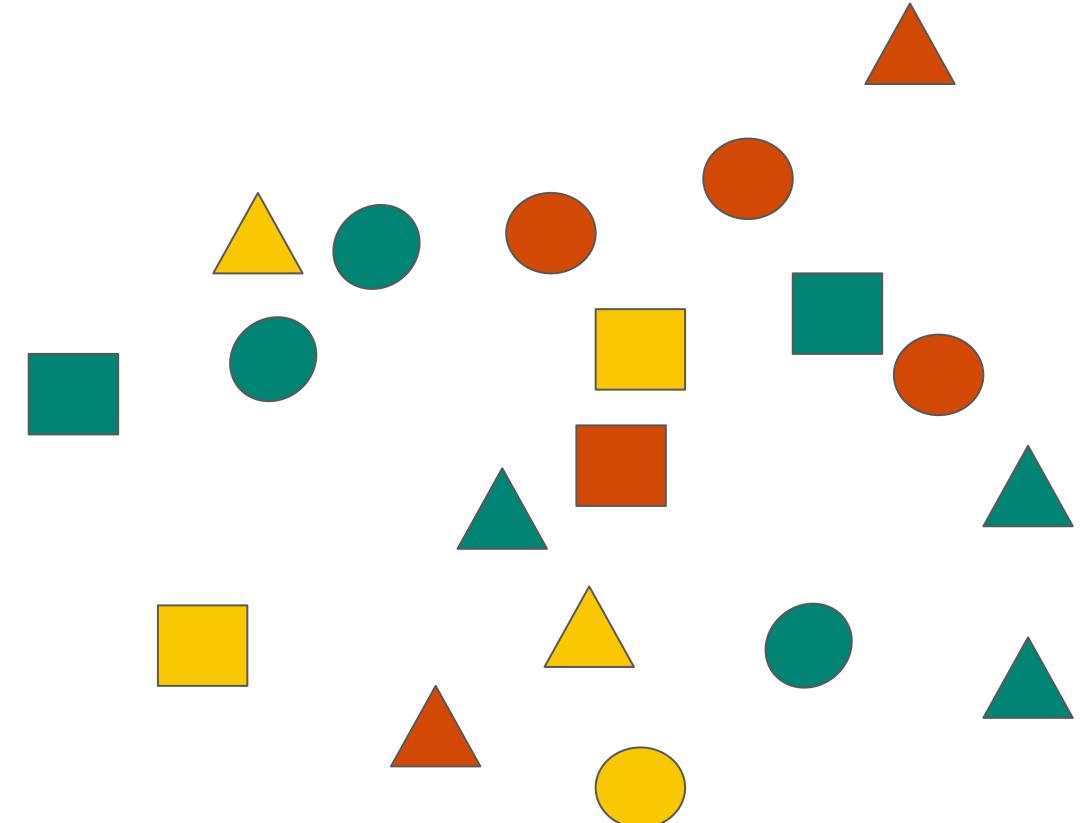
Artificial Intelligence vs. Machine Learning

Machines that mimic human cognition

Building machines that can solve problems and make decisions

A form of artificial intelligence

Pattern matching. It “learns” based on training data.



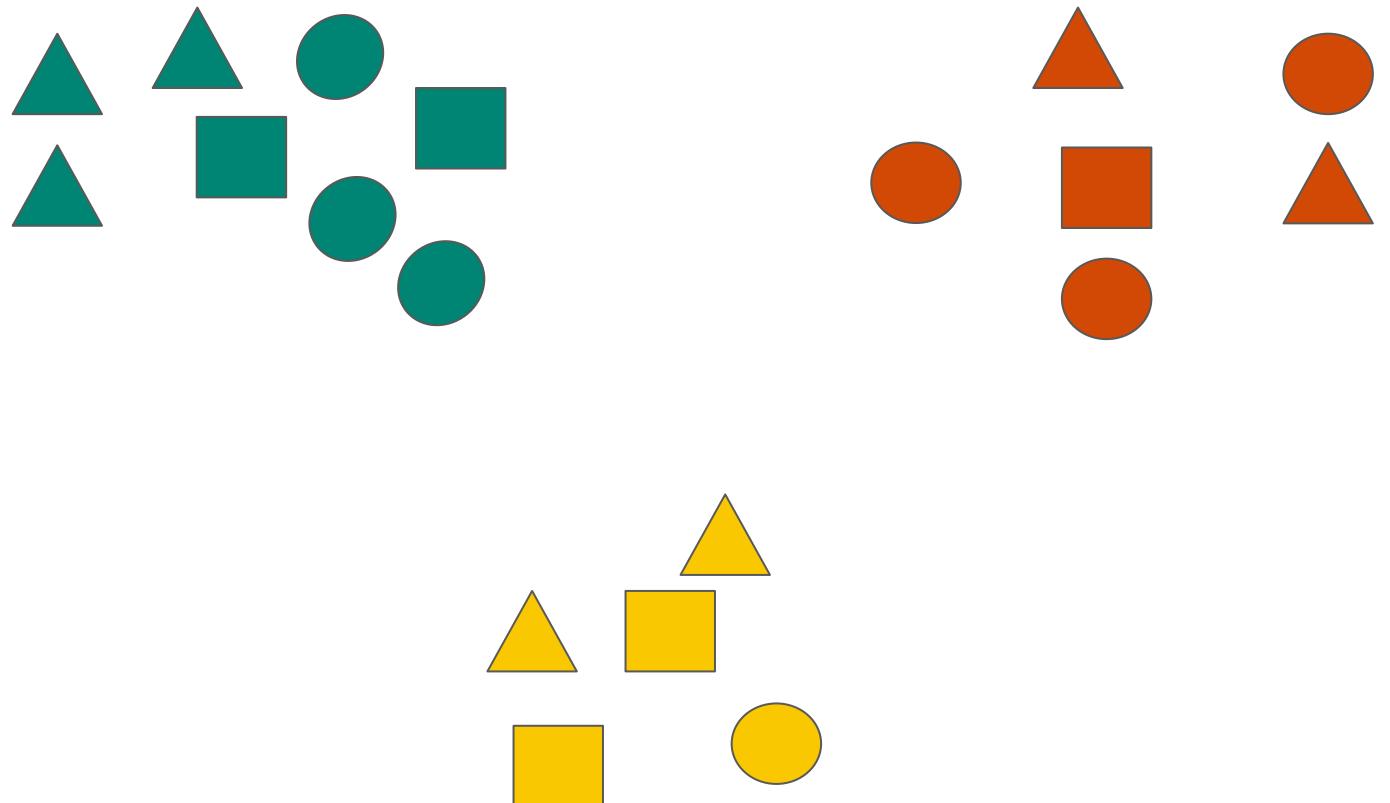
Artificial Intelligence vs. Machine Learning

Machines that mimic human cognition

Building machines that can solve problems and make decisions

A form of artificial intelligence

Pattern matching. It “learns” based on training data.



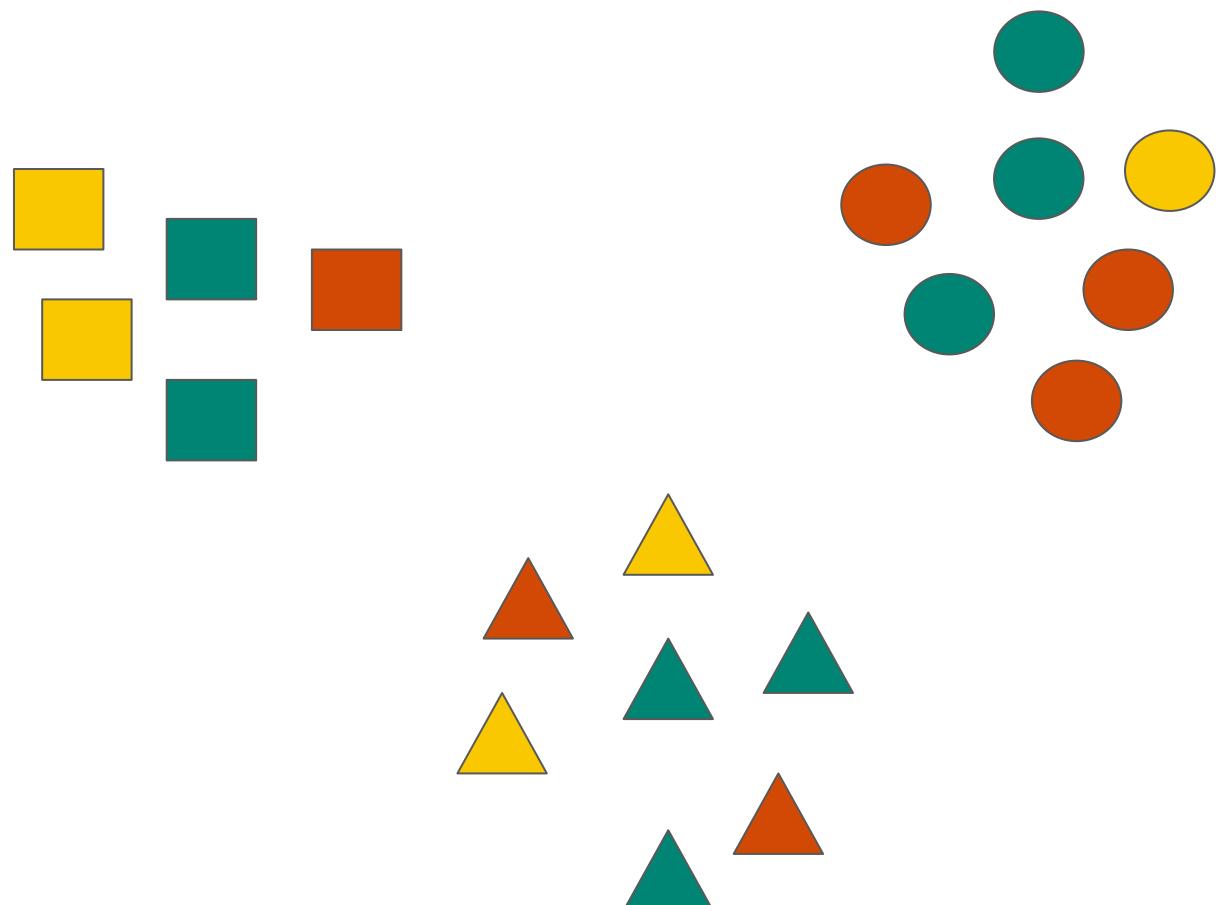
Artificial Intelligence vs. Machine Learning

Machines that mimic human cognition

Building machines that can solve problems and make decisions

A form of artificial intelligence

Pattern matching. It “learns” based on training data.



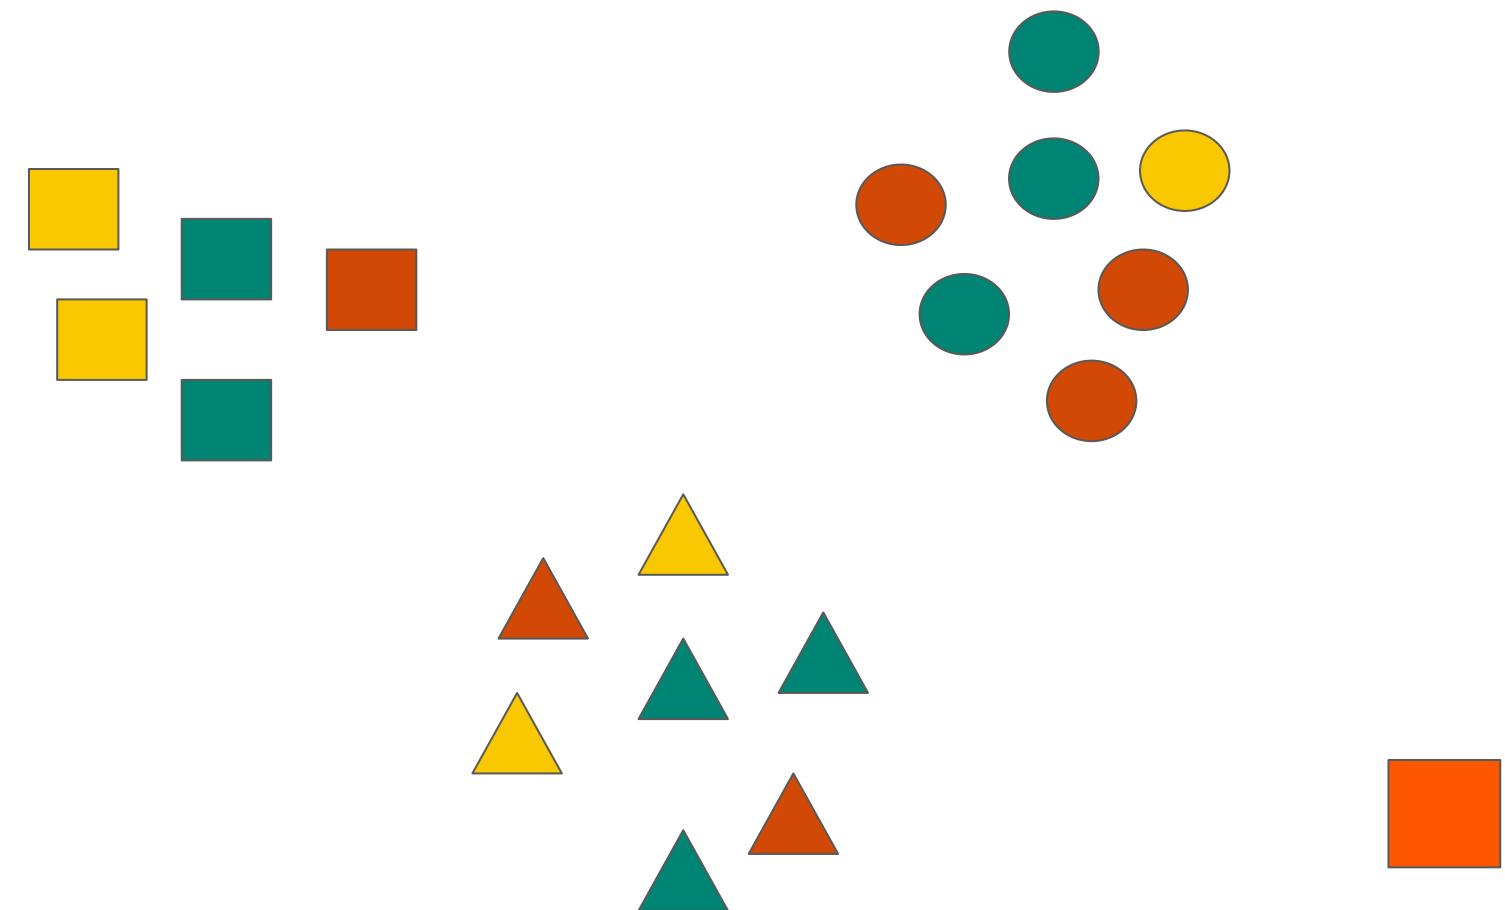
Artificial Intelligence vs. Machine Learning

Machines that mimic human cognition

Building machines that can solve problems and make decisions

A form of artificial intelligence

Pattern matching. It “learns” based on training data.



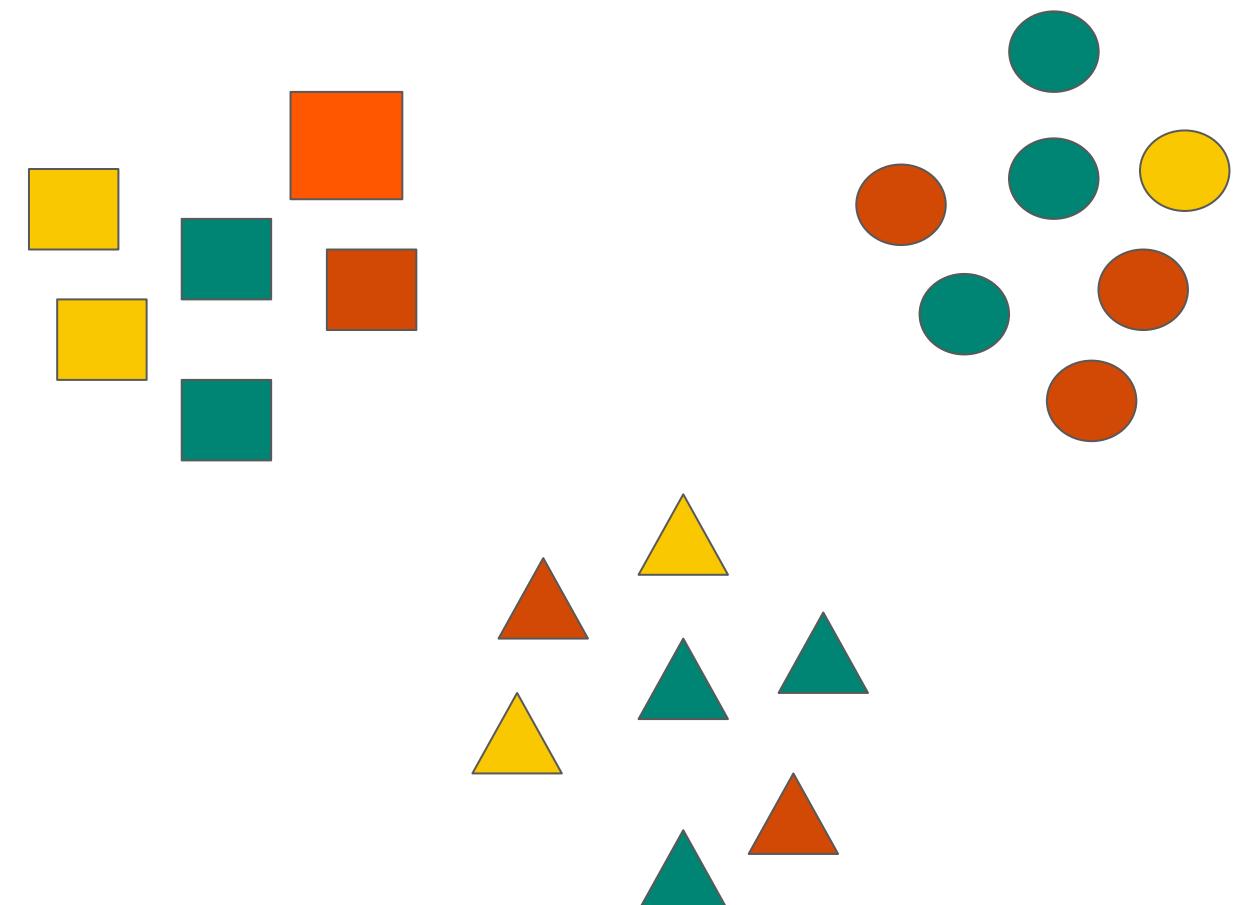
Artificial Intelligence vs. Machine Learning

Machines that mimic human cognition

Building machines that can solve problems and make decisions

A form of artificial intelligence

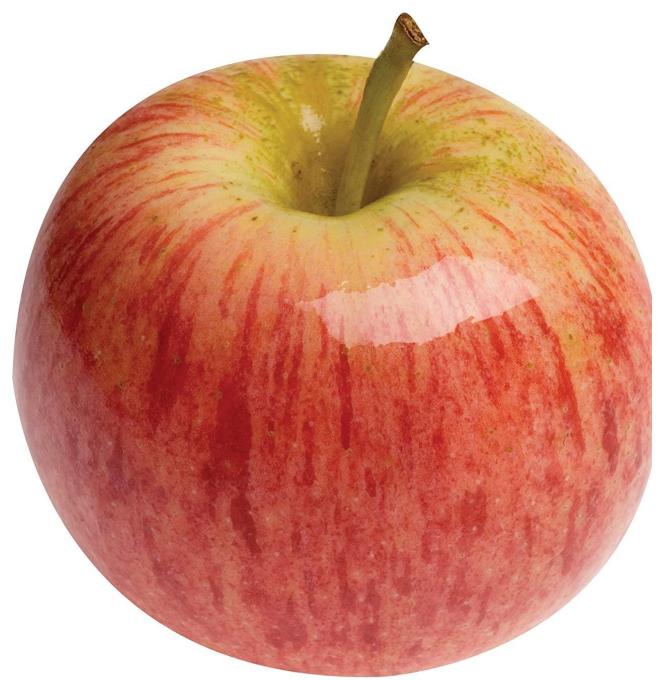
Pattern matching. It “learns” based on training data.





Bias

N



Apple



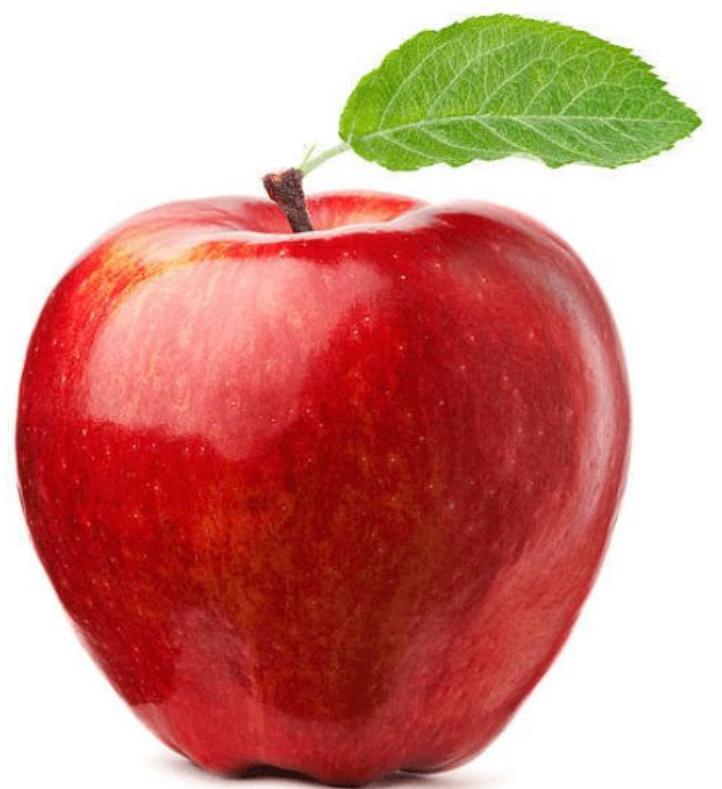
Banana



Apple



Apple



Apple



Apple

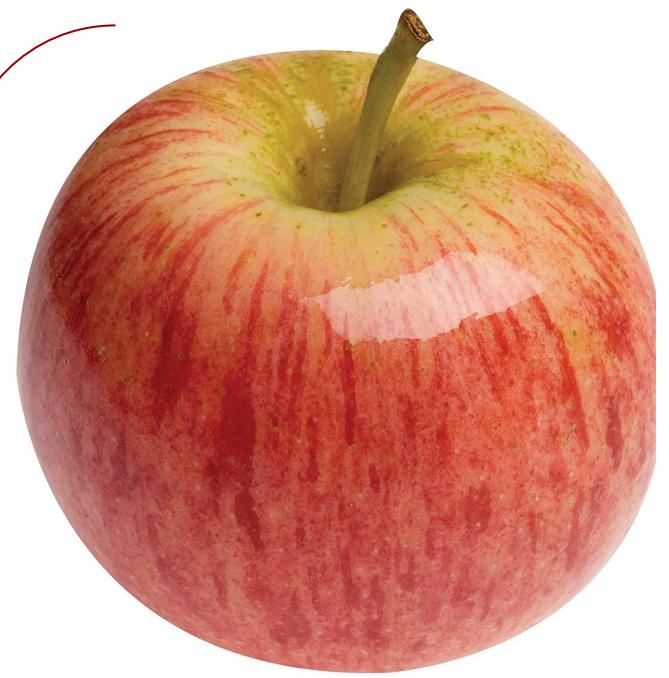


Banana



Apple

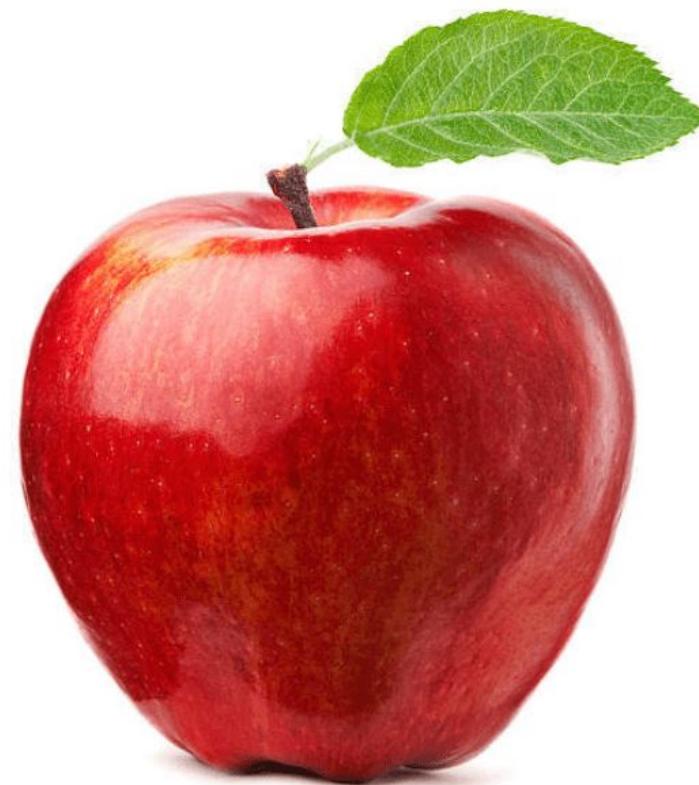
N



Apple



Banana



Apple



Apple

Learn a model from
this training data

Model

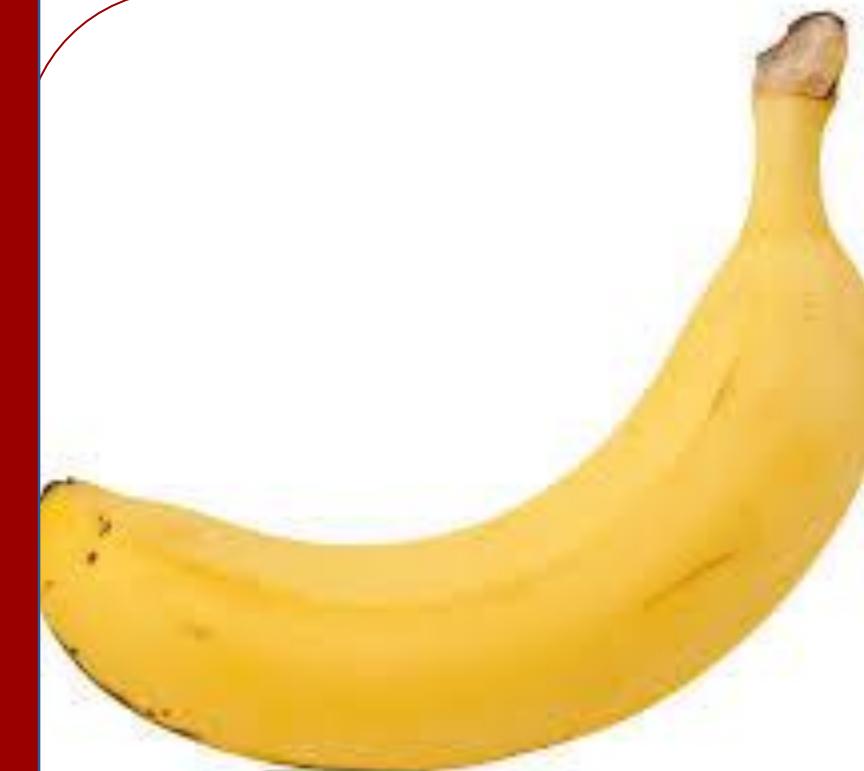
Model



Apple



Apple

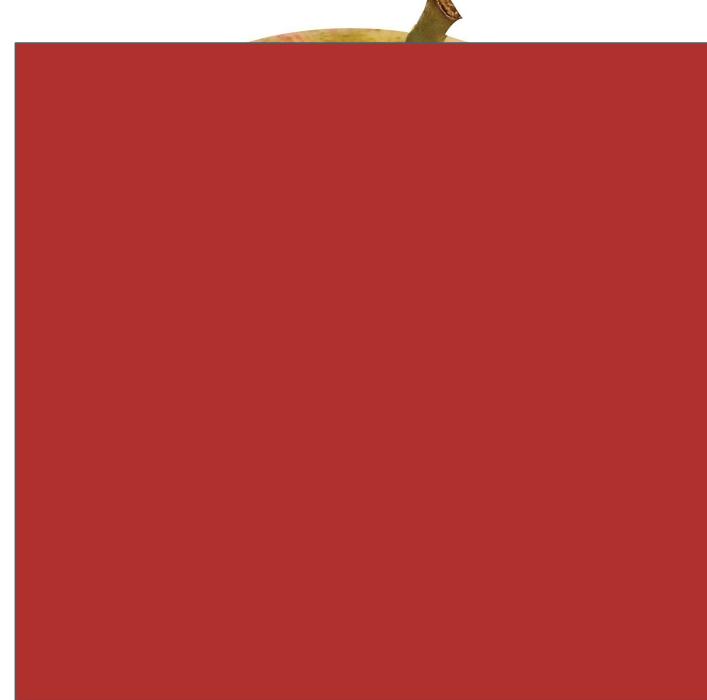


Banana



Banana

Why?



Apple



Banana

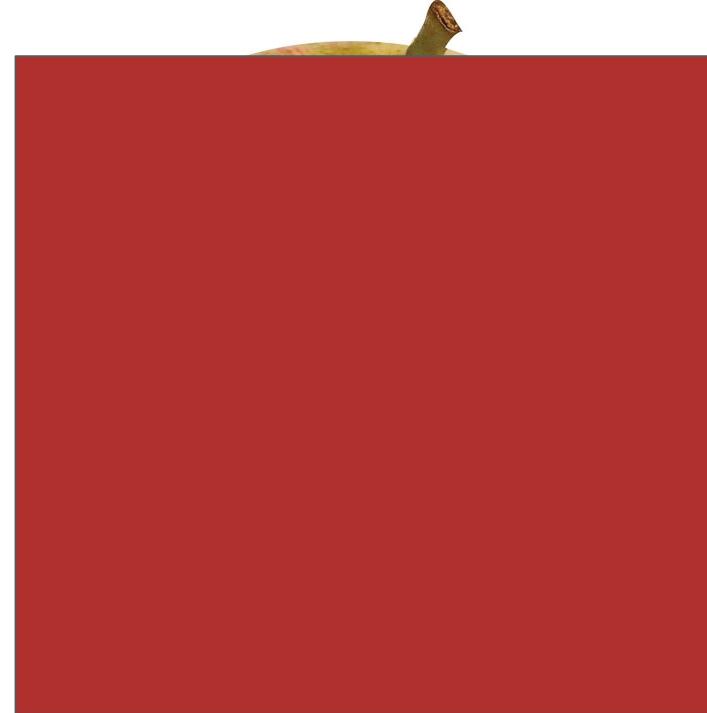


Apple



Apple

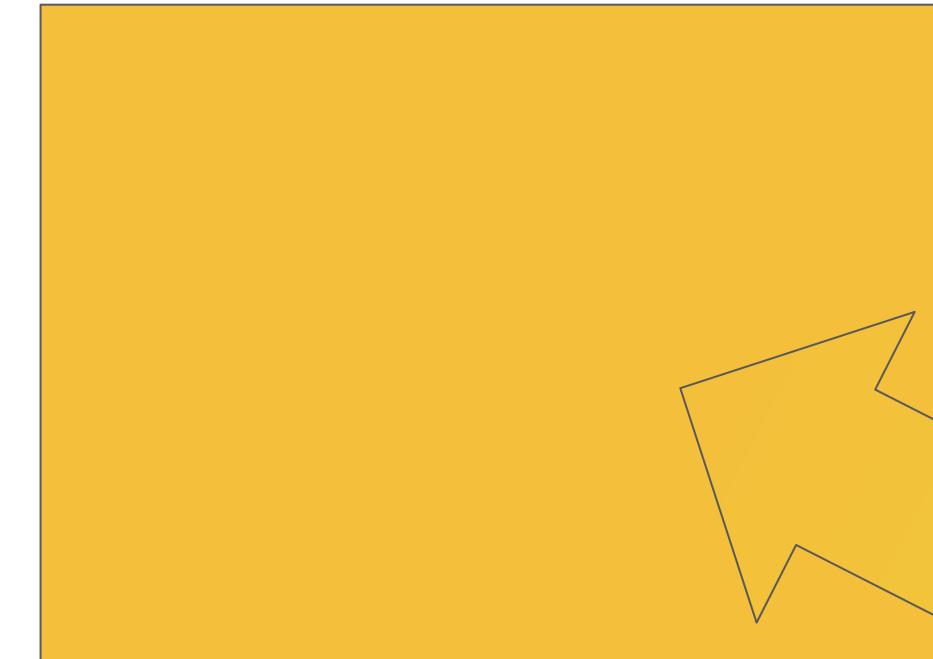
Why?



Apple



Apple

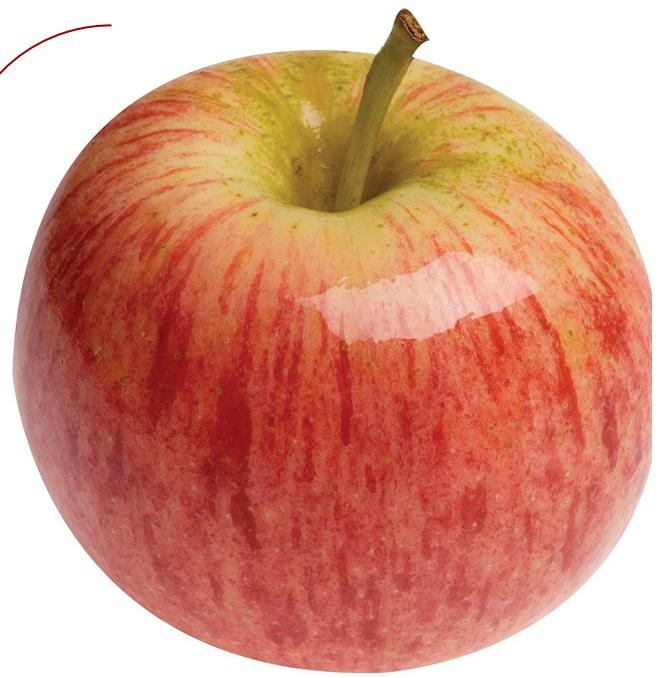


Banana



Apple

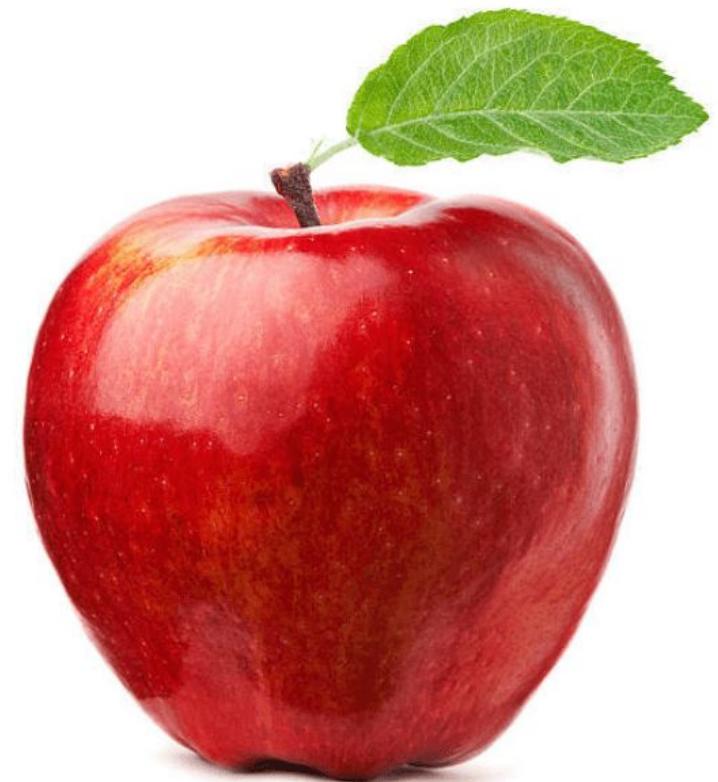
N



Apple



Banana

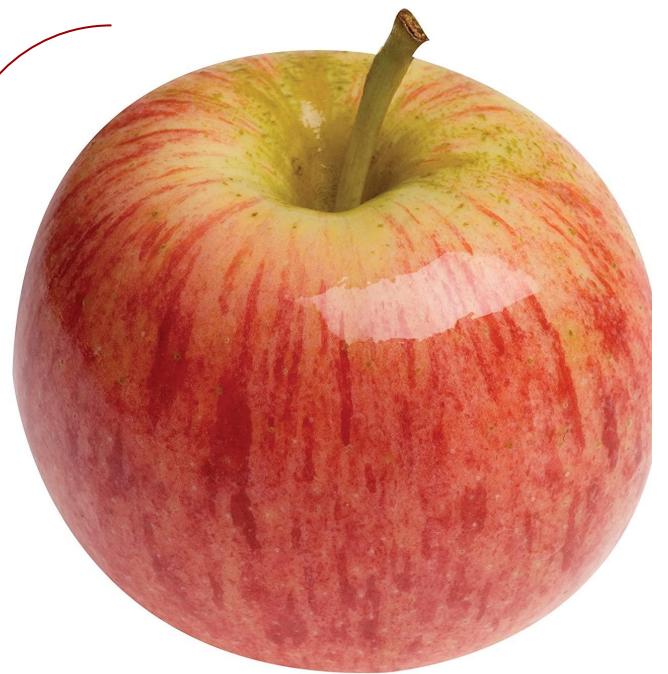


Apple



Apple

N



Apple



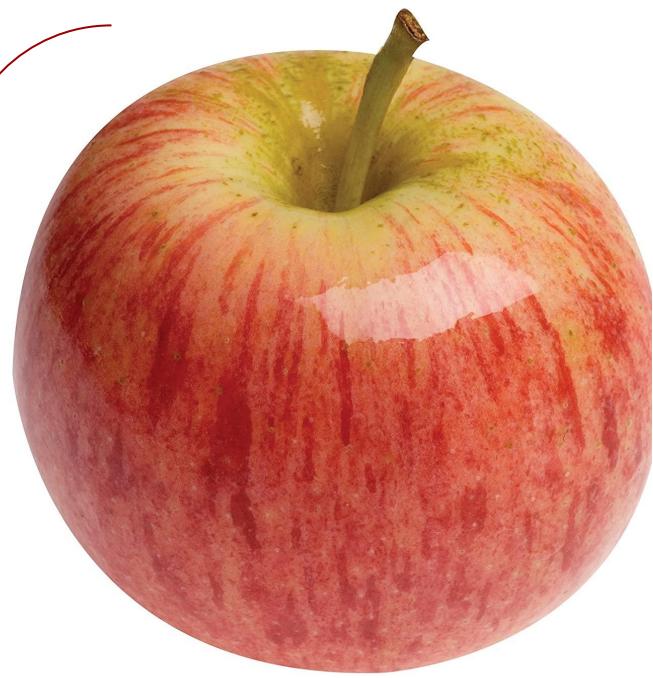
Banana



Apple

Apple

N



Apple



Banana



Apple

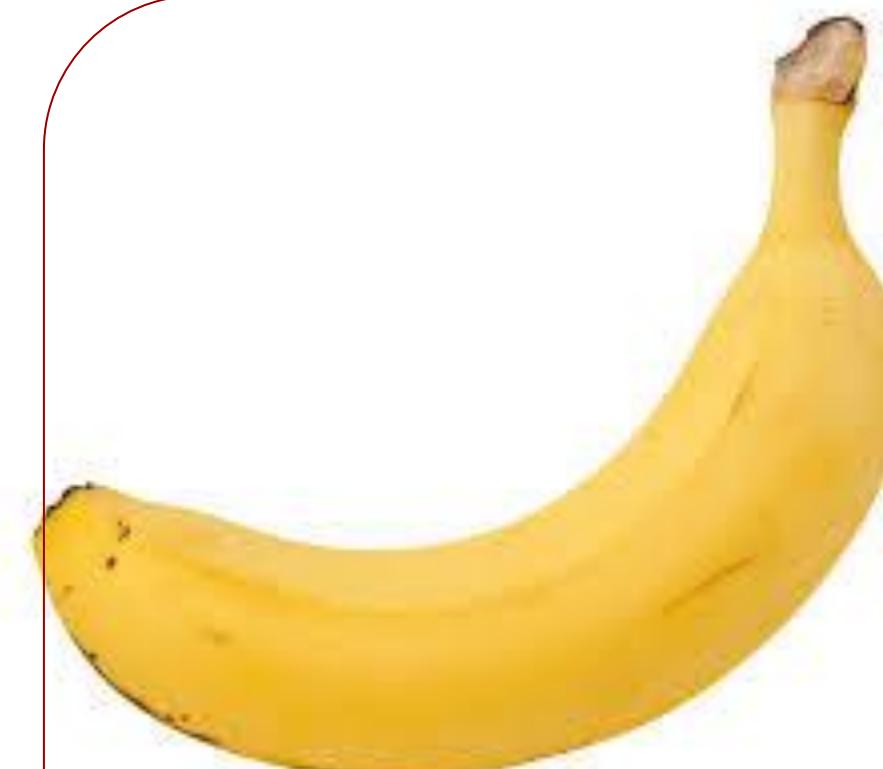
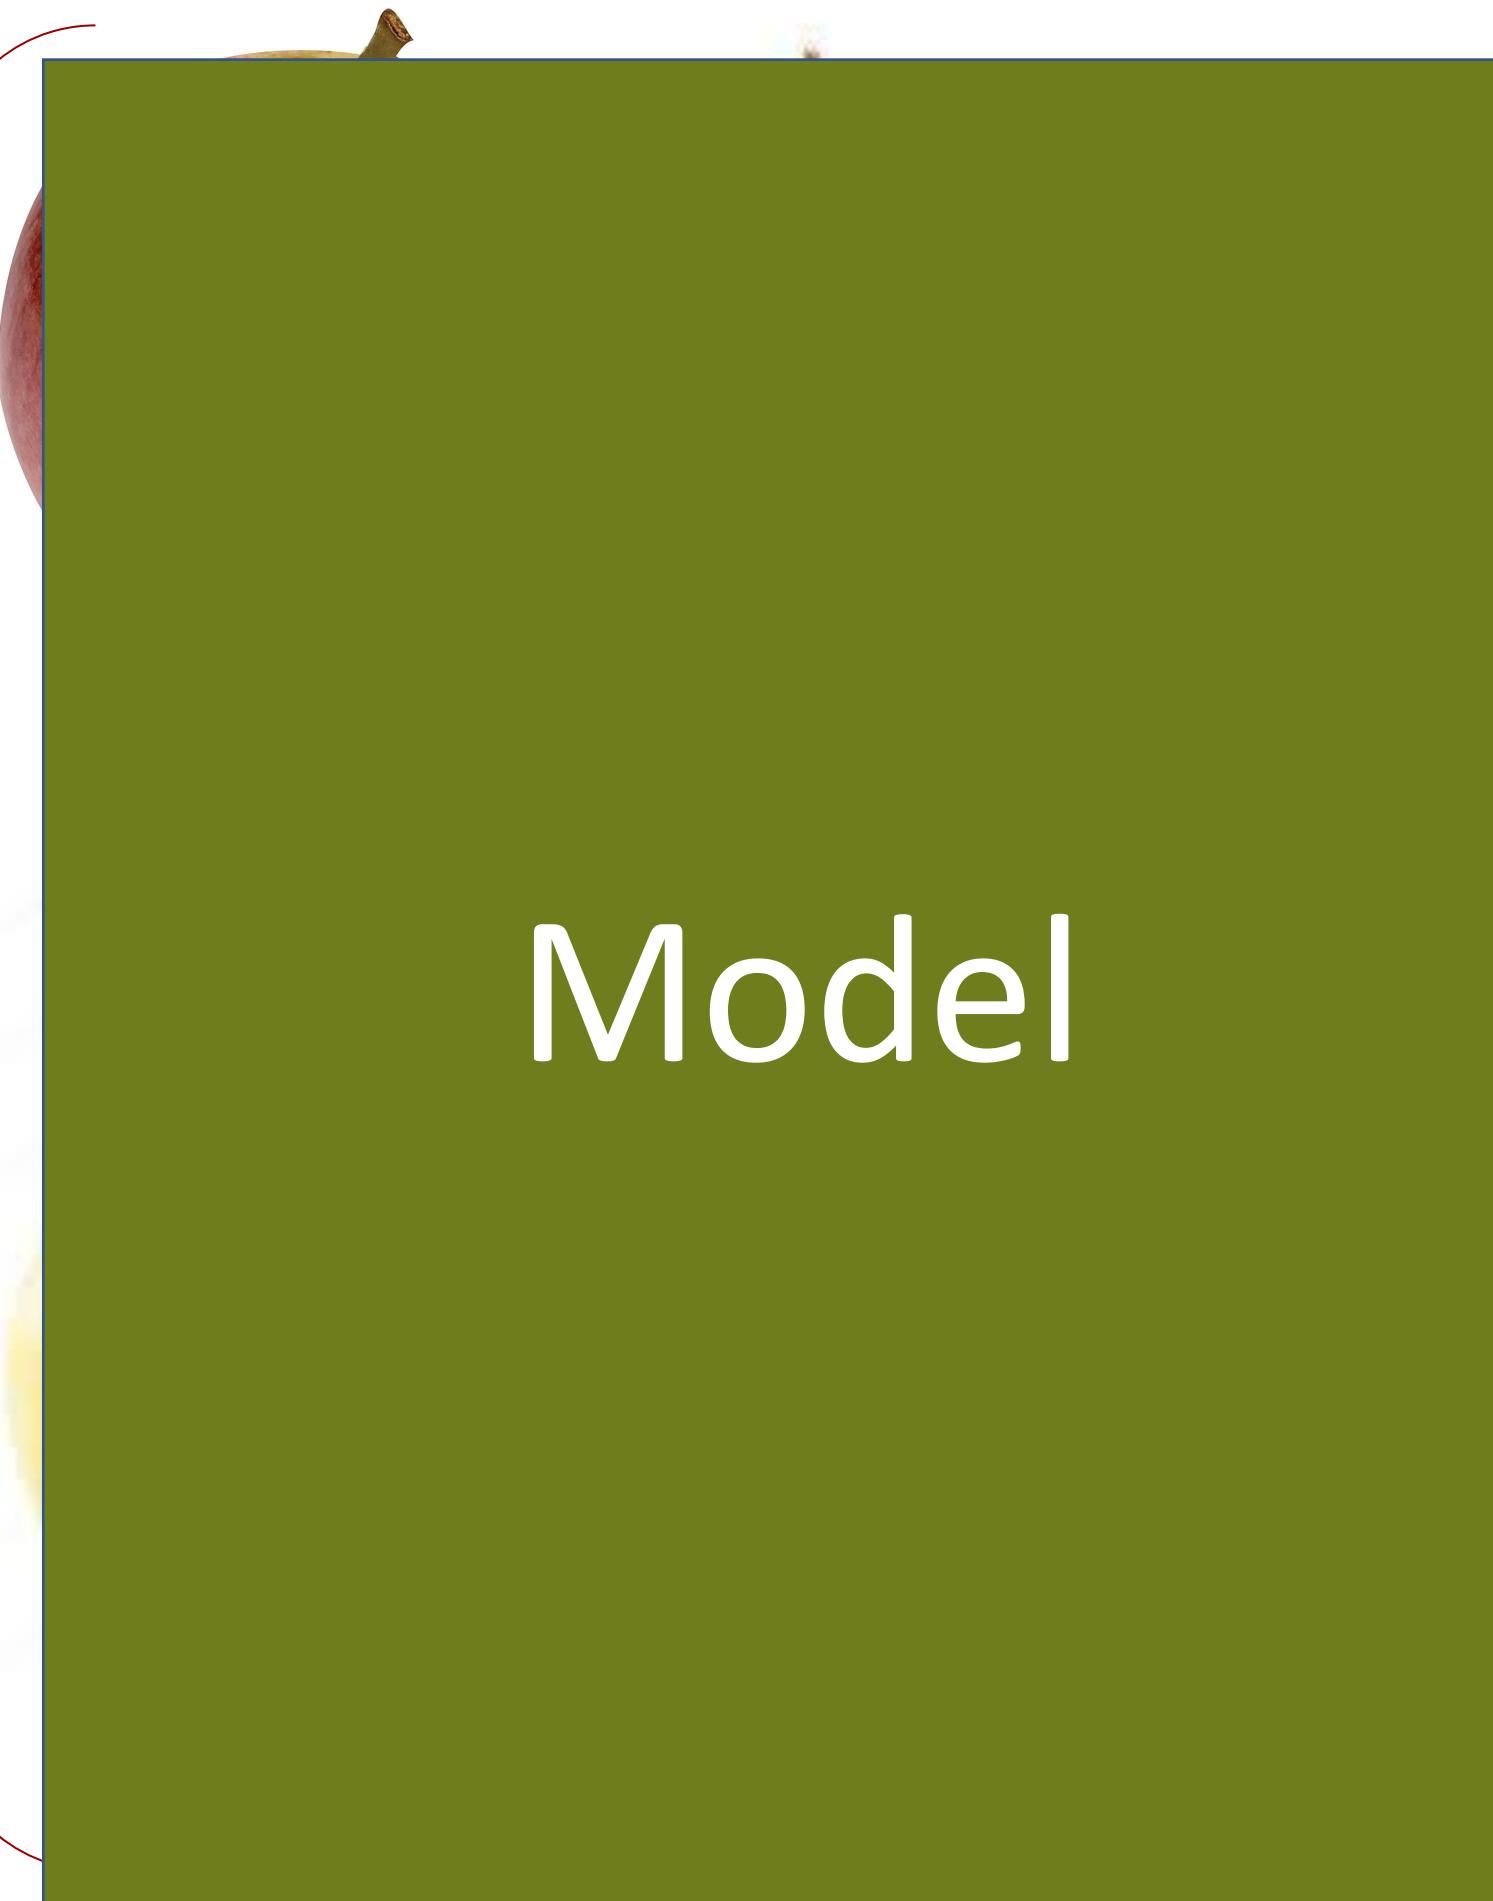


Apple

Learn a model from
this more diverse
training data

N

Model



Apple



Apple



Banana

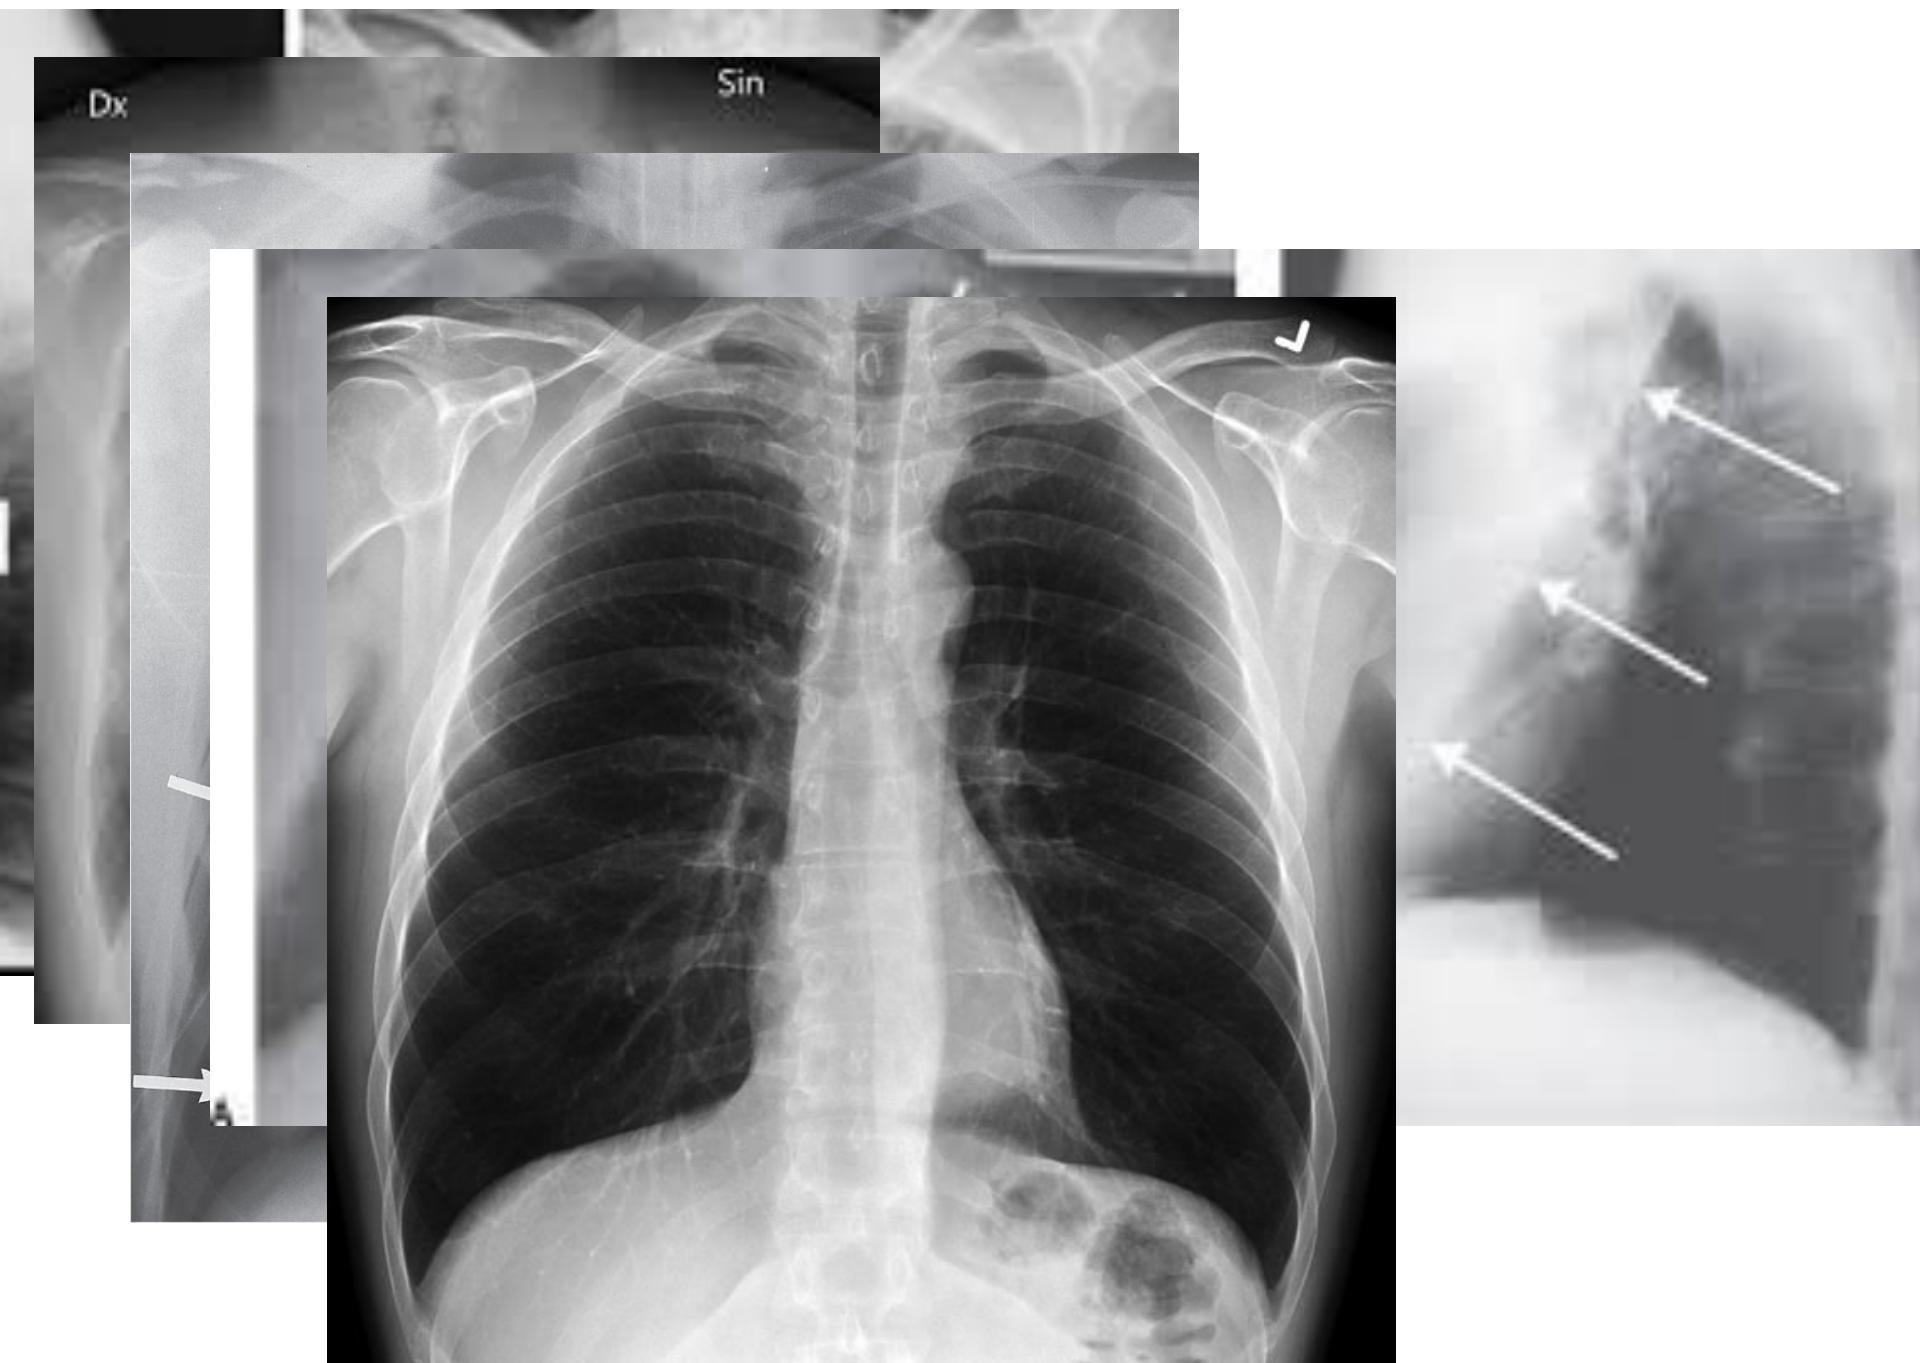


Apple



Let's make this more
real-world

Imagine you are a radiologist.



Cancer

Healthy

Images



Cancer

Healthy

Images

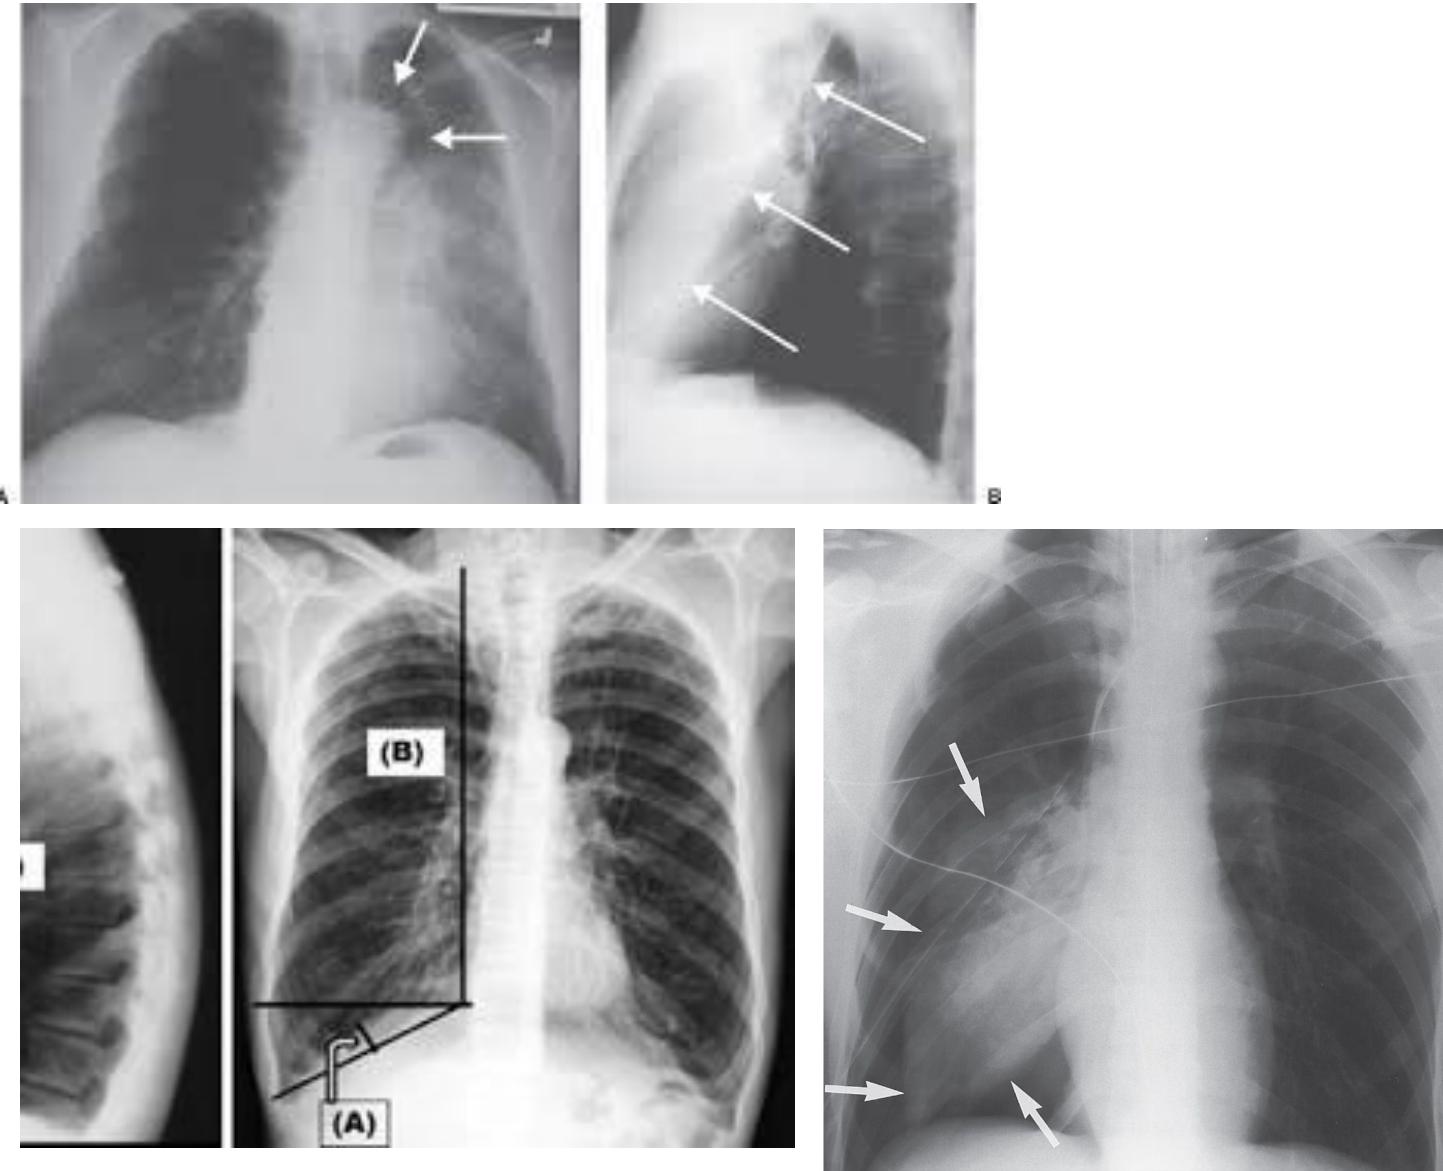
Machine
Learning



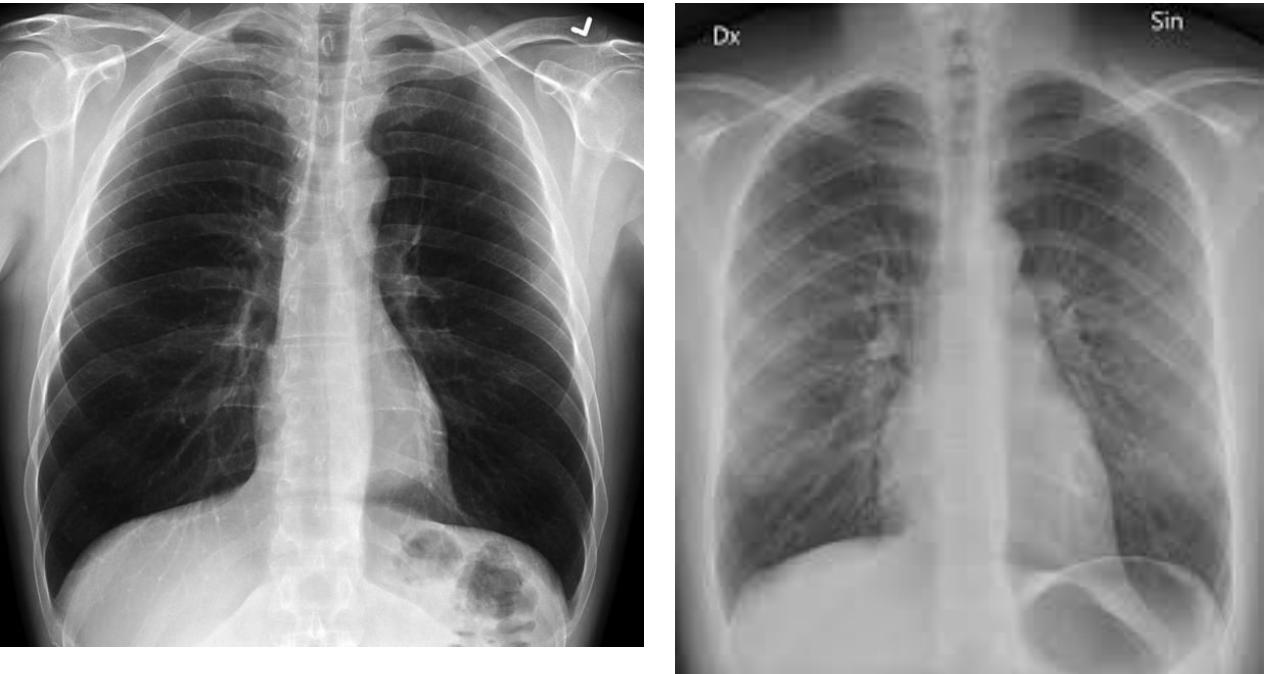
Cancer

Healthy

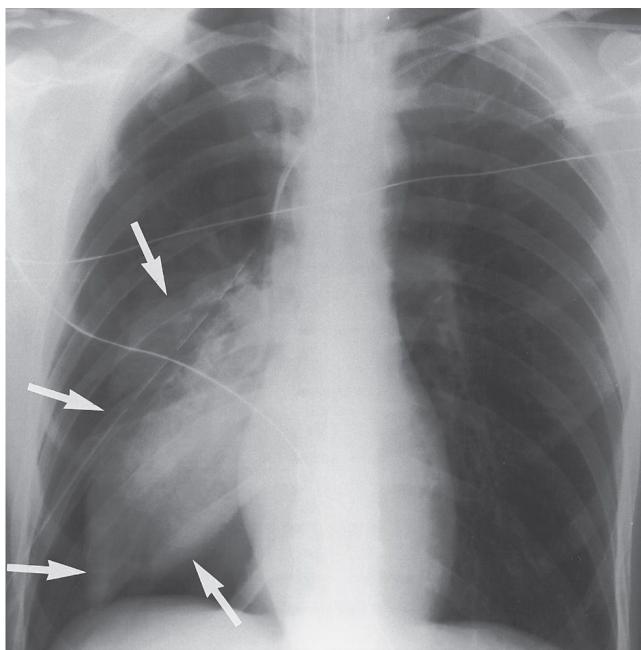
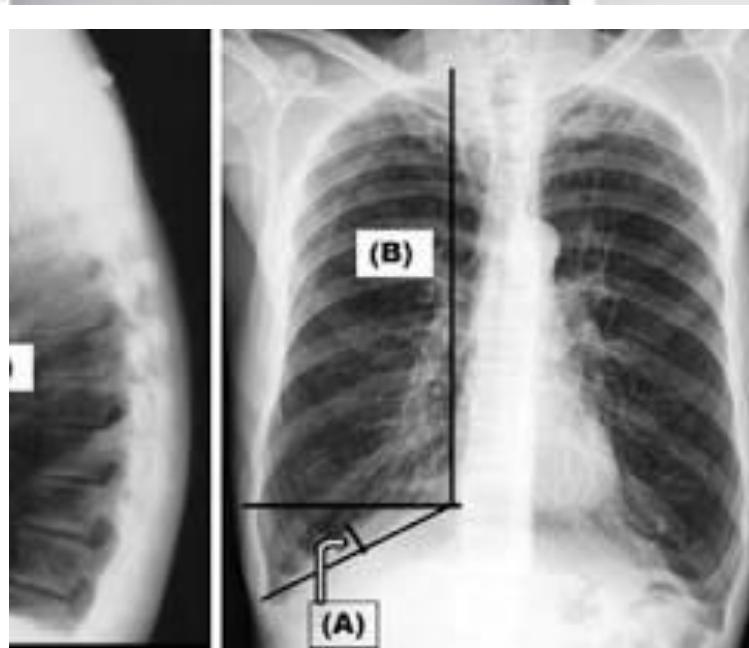
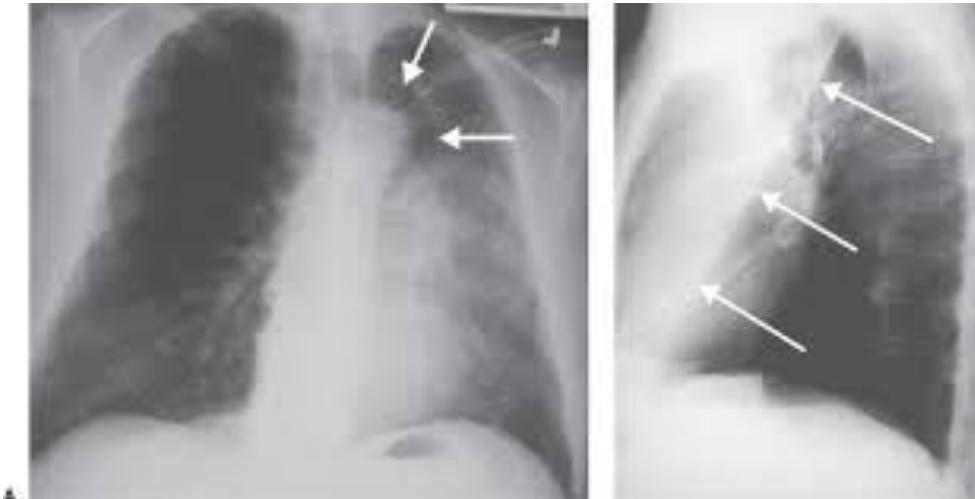
Cancer



Healthy



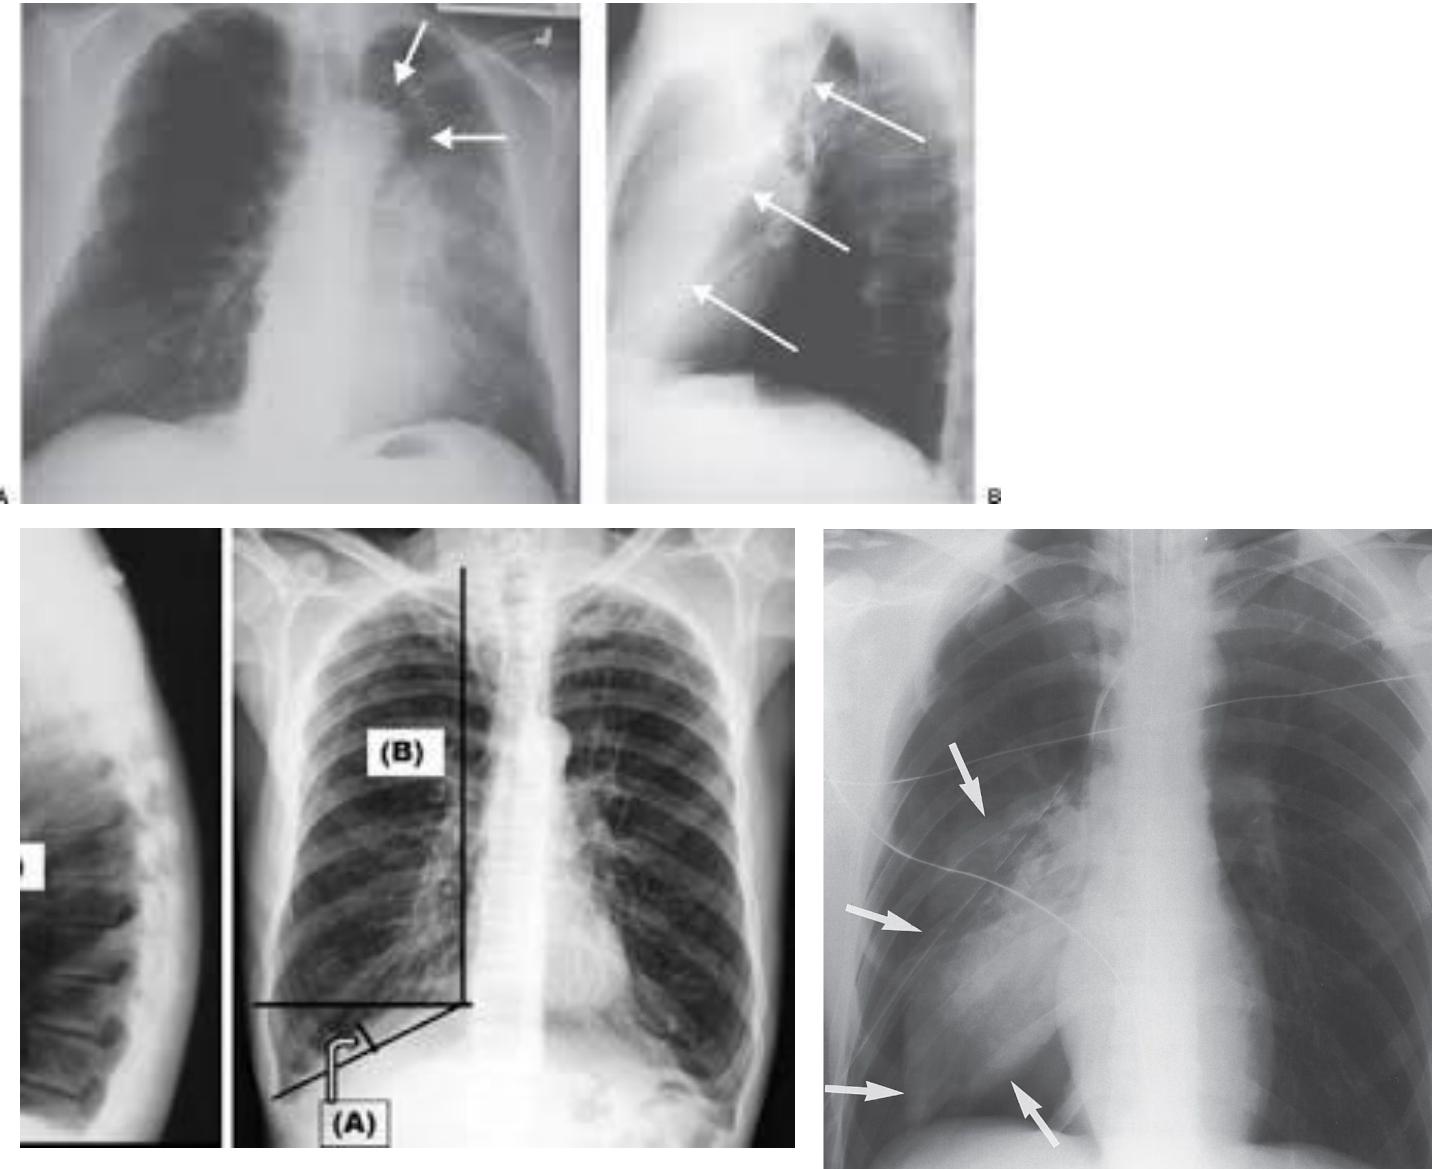
Cancer



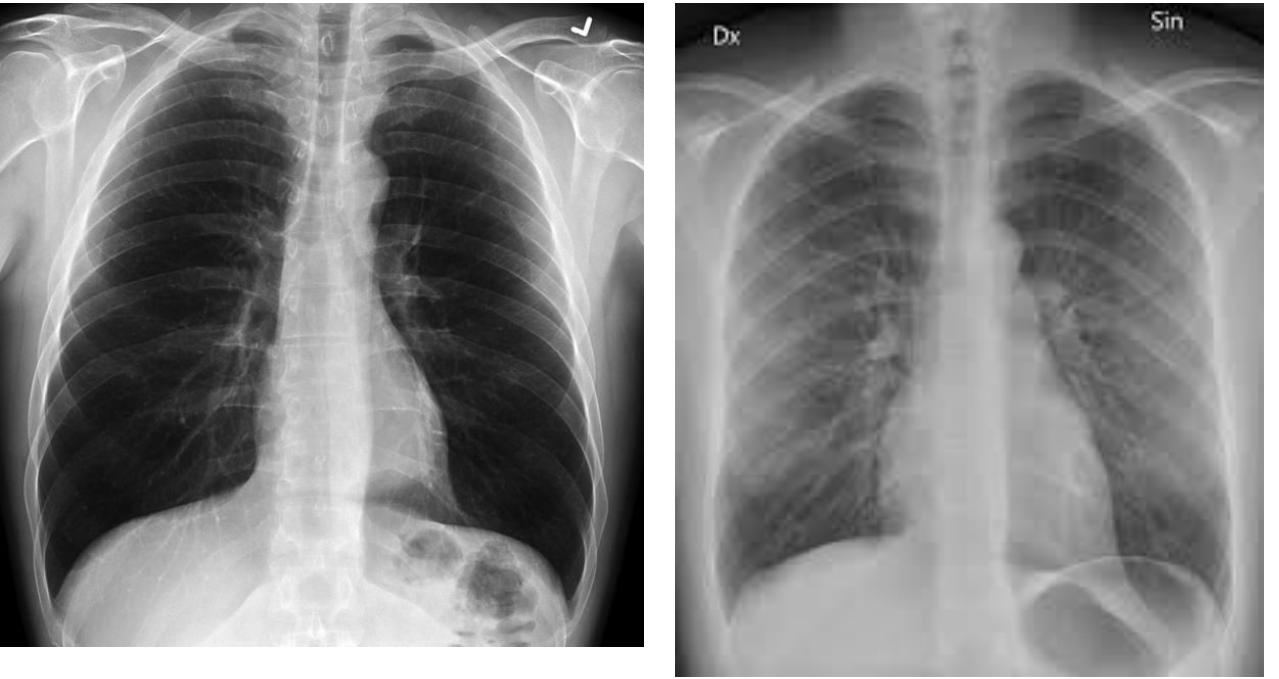
Healthy



Cancer



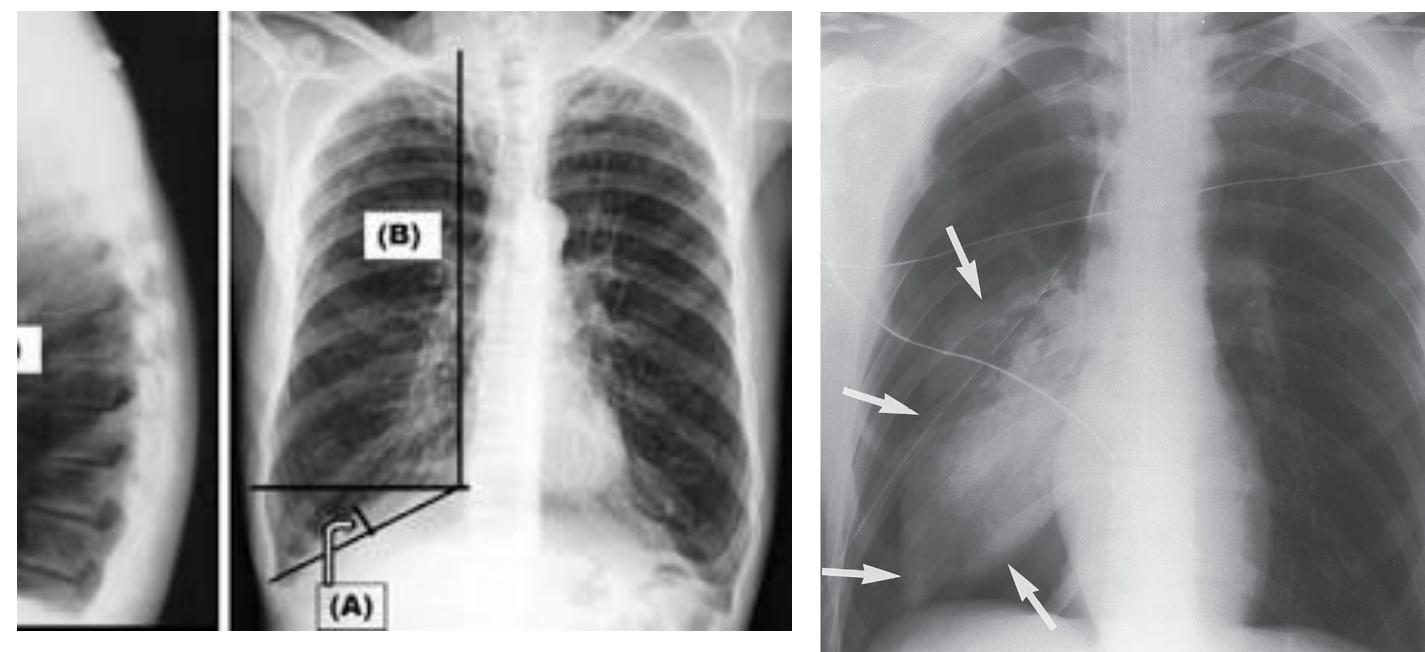
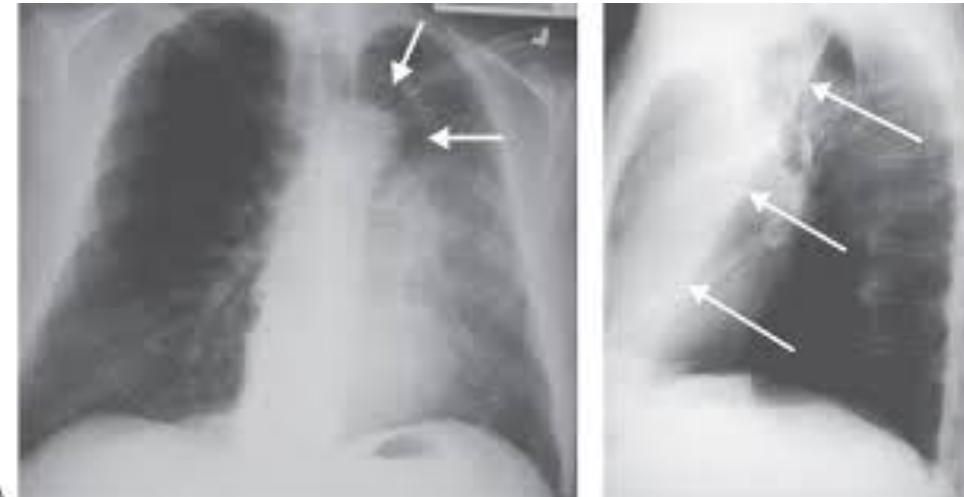
Healthy



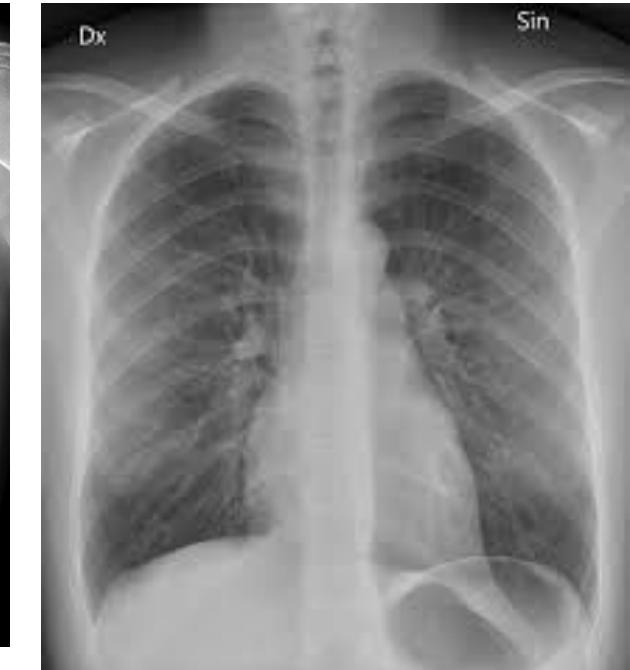
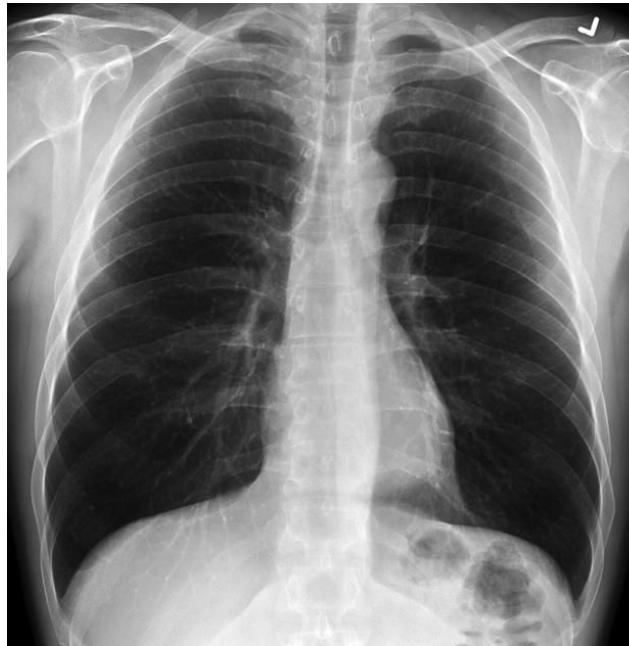
!!



Cancer



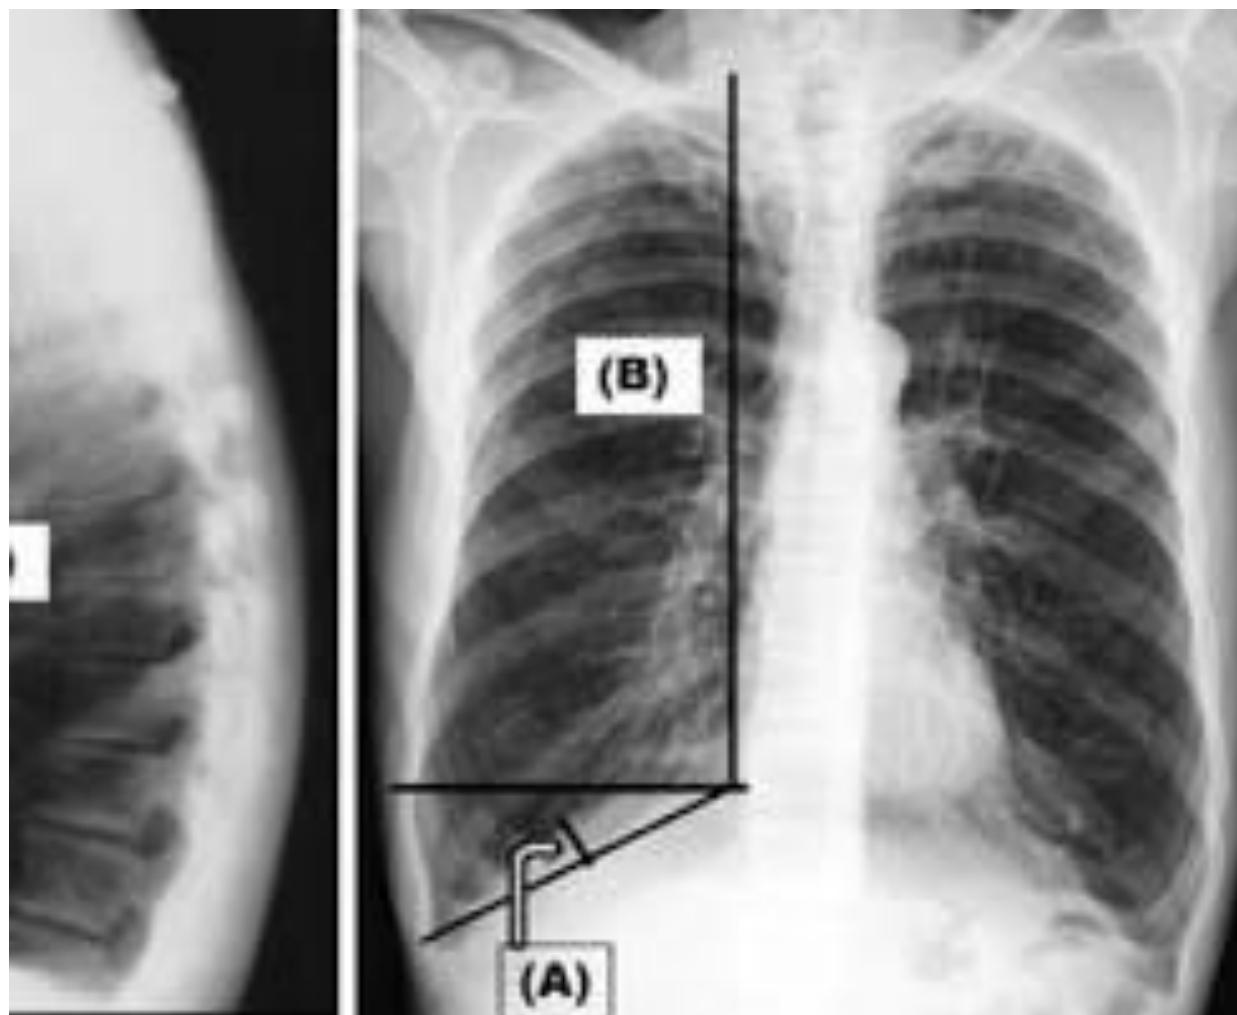
Healthy



!!



This was real.



VB Live

When AI flags the ruler, not the tumor – and other arguments for abolishing the black box (VB Live)

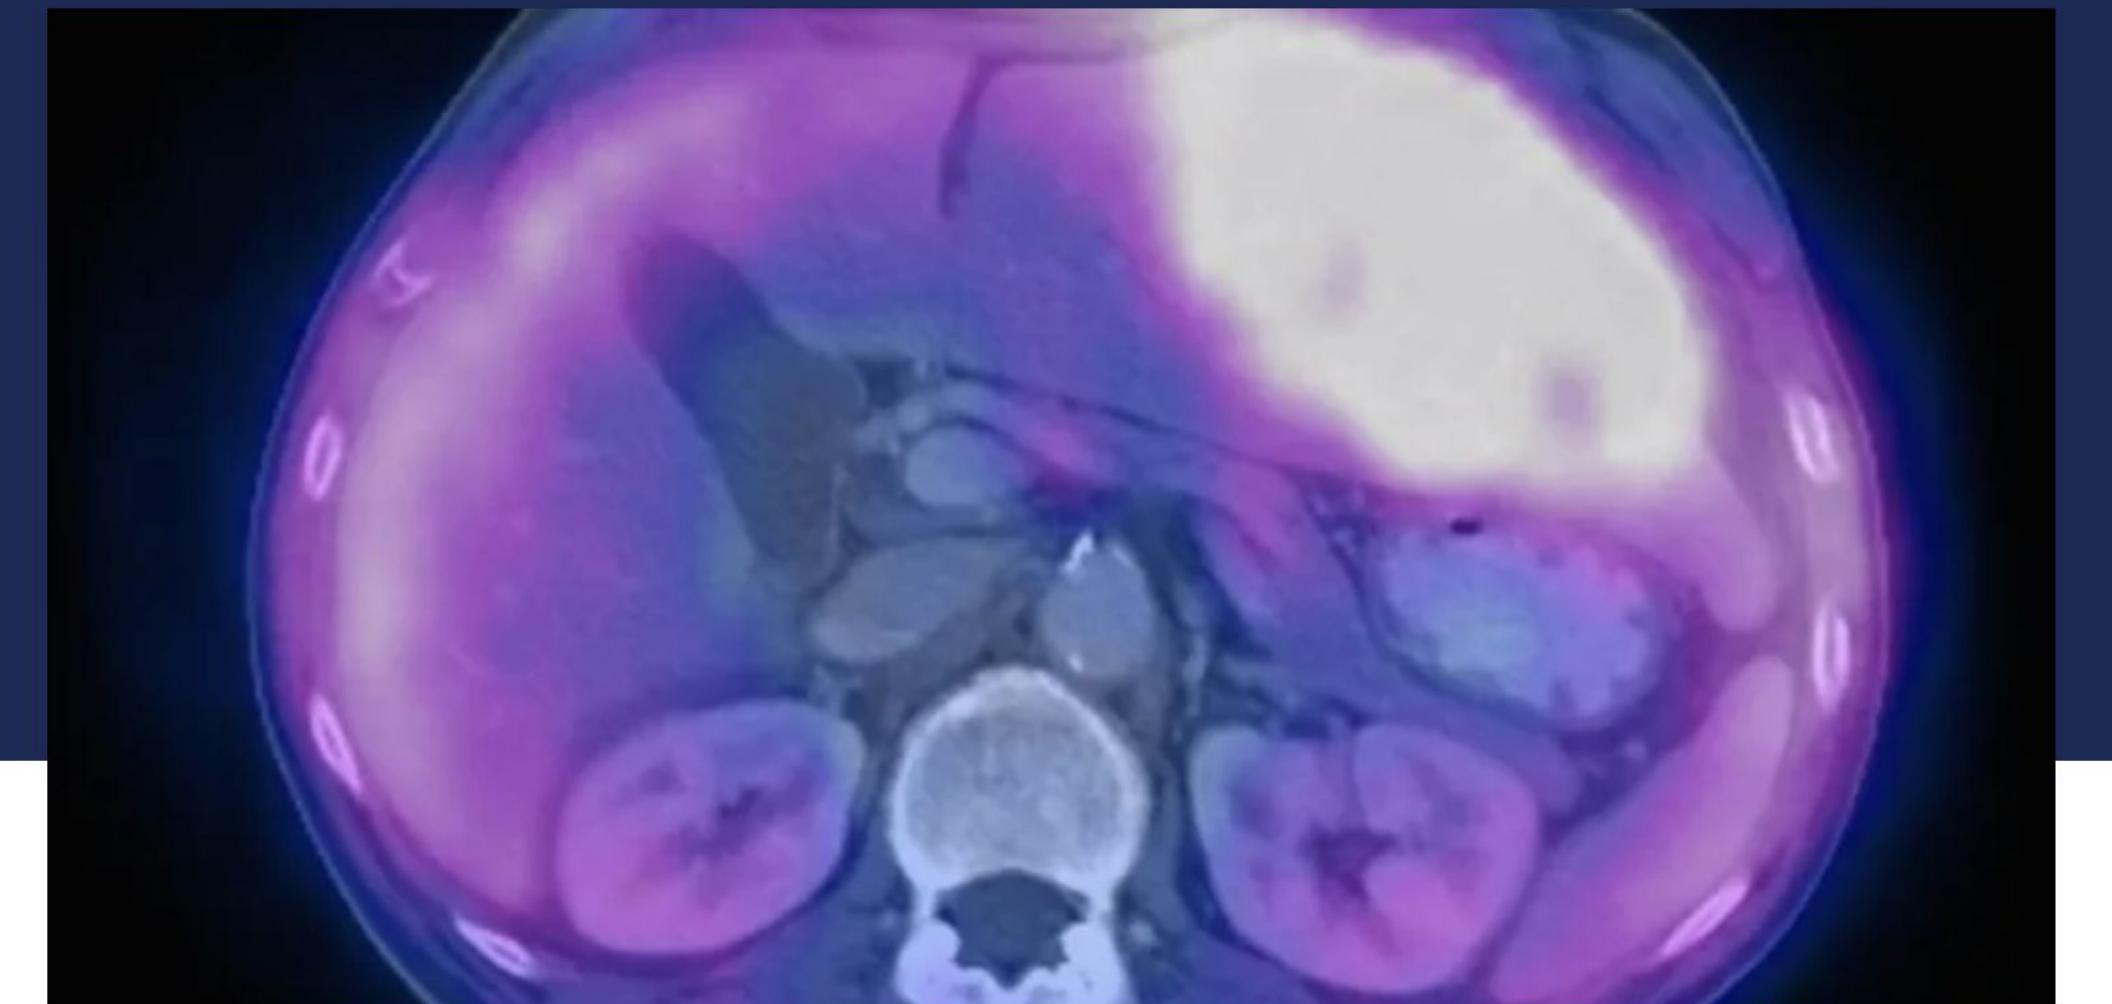


Image Credit: Getty Images

:he

AI helps health care experts do their jobs efficiently and effectively, but it needs to be used responsibly, ethically, and equitably. In this VB Live event, get an in-depth perspective on the strengths and limitations of data, AI methodology and more.

VB
Ma
f

What is ChatGPT?

Start simple: A bigram language model

Let's build a model based on counting. This is where prediction began.

We care about predicting *the next word* after a single word.

This is called a **bigram**.

From fairest creatures we desire increase,
That thereby beauty's rose might never die,
But as the riper should by time decease,
His tender heir might bear his memory;

	from	fairest	creatures	we	desire
from					
fairest					
creatures					
we					
desire					

From fairest creatures we desire increase,
That thereby beauty's rose might never die,
But as the riper should by time decease,
His tender heir might bear his memory;

	from	fairest	creatures	we	youth
from	0	0	0	0	0
fairest	0	0	0	0	0
creatures	0	0	0	0	0
we	0	0	0	0	0
youth	0	0	0	0	0

From fairest creatures we desire increase,
That thereby beauty's rose might never die,
But as the riper should by time decease,
His tender heir might bear his memory;

	from	fairest	creatures	we	youth
from	0	1	0	0	0
fairest	0	0	0	0	0
creatures	0	0	0	0	0
we	0	0	0	0	0
youth	0	0	0	0	0

From fairest creatures we desire increase,
That thereby beauty's rose might never die,
But as the riper should by time decease,
His tender heir might bear his memory;

	from	fairest	creatures	we	youth
from	0	1	0	0	0
fairest	0	0	1	0	0
creatures	0	0	0	0	0
we	0	0	0	0	0
youth	0	0	0	0	0

From fairest creatures we desire increase,
 That thereby beauty's rose might never die,
 But as the riper should by time decease,
 His tender heir might bear his memory;

	from	fairest	creatures	we	youth
from	0	1	0	0	0
fairest	0	0	1	0	0
creatures	0	0	0	1	0
we	0	0	0	0	0
youth	0	0	0	0	0

From fairest creatures we desire increase,
 That thereby beauty's rose might never die,
 But as the riper should by time decease,
 His tender heir might bear his memory;

.

.

.

From youth convertest.

	from	fairest	creatures	we	youth
from	0	1	0	0	0
fairest	0	0	1	0	0
creatures	0	0	0	1	0
we	0	0	0	0	0
youth	0	0	0	0	0

And so on and so forth for the
complete works of
Shakespeare.

From fairest creatures we

That thereby beauty's rose ~~might never die,~~

But as the riper should by time decease,

His tender heir might bear his memory;

.

.

.

From youth convertest.

	airest	creatures	we	youth
fairest	0	0	1	0
creatures	0	0	0	1
we	0	0	0	0
youth	0	0	0	0

Inference

Now our model is trained.

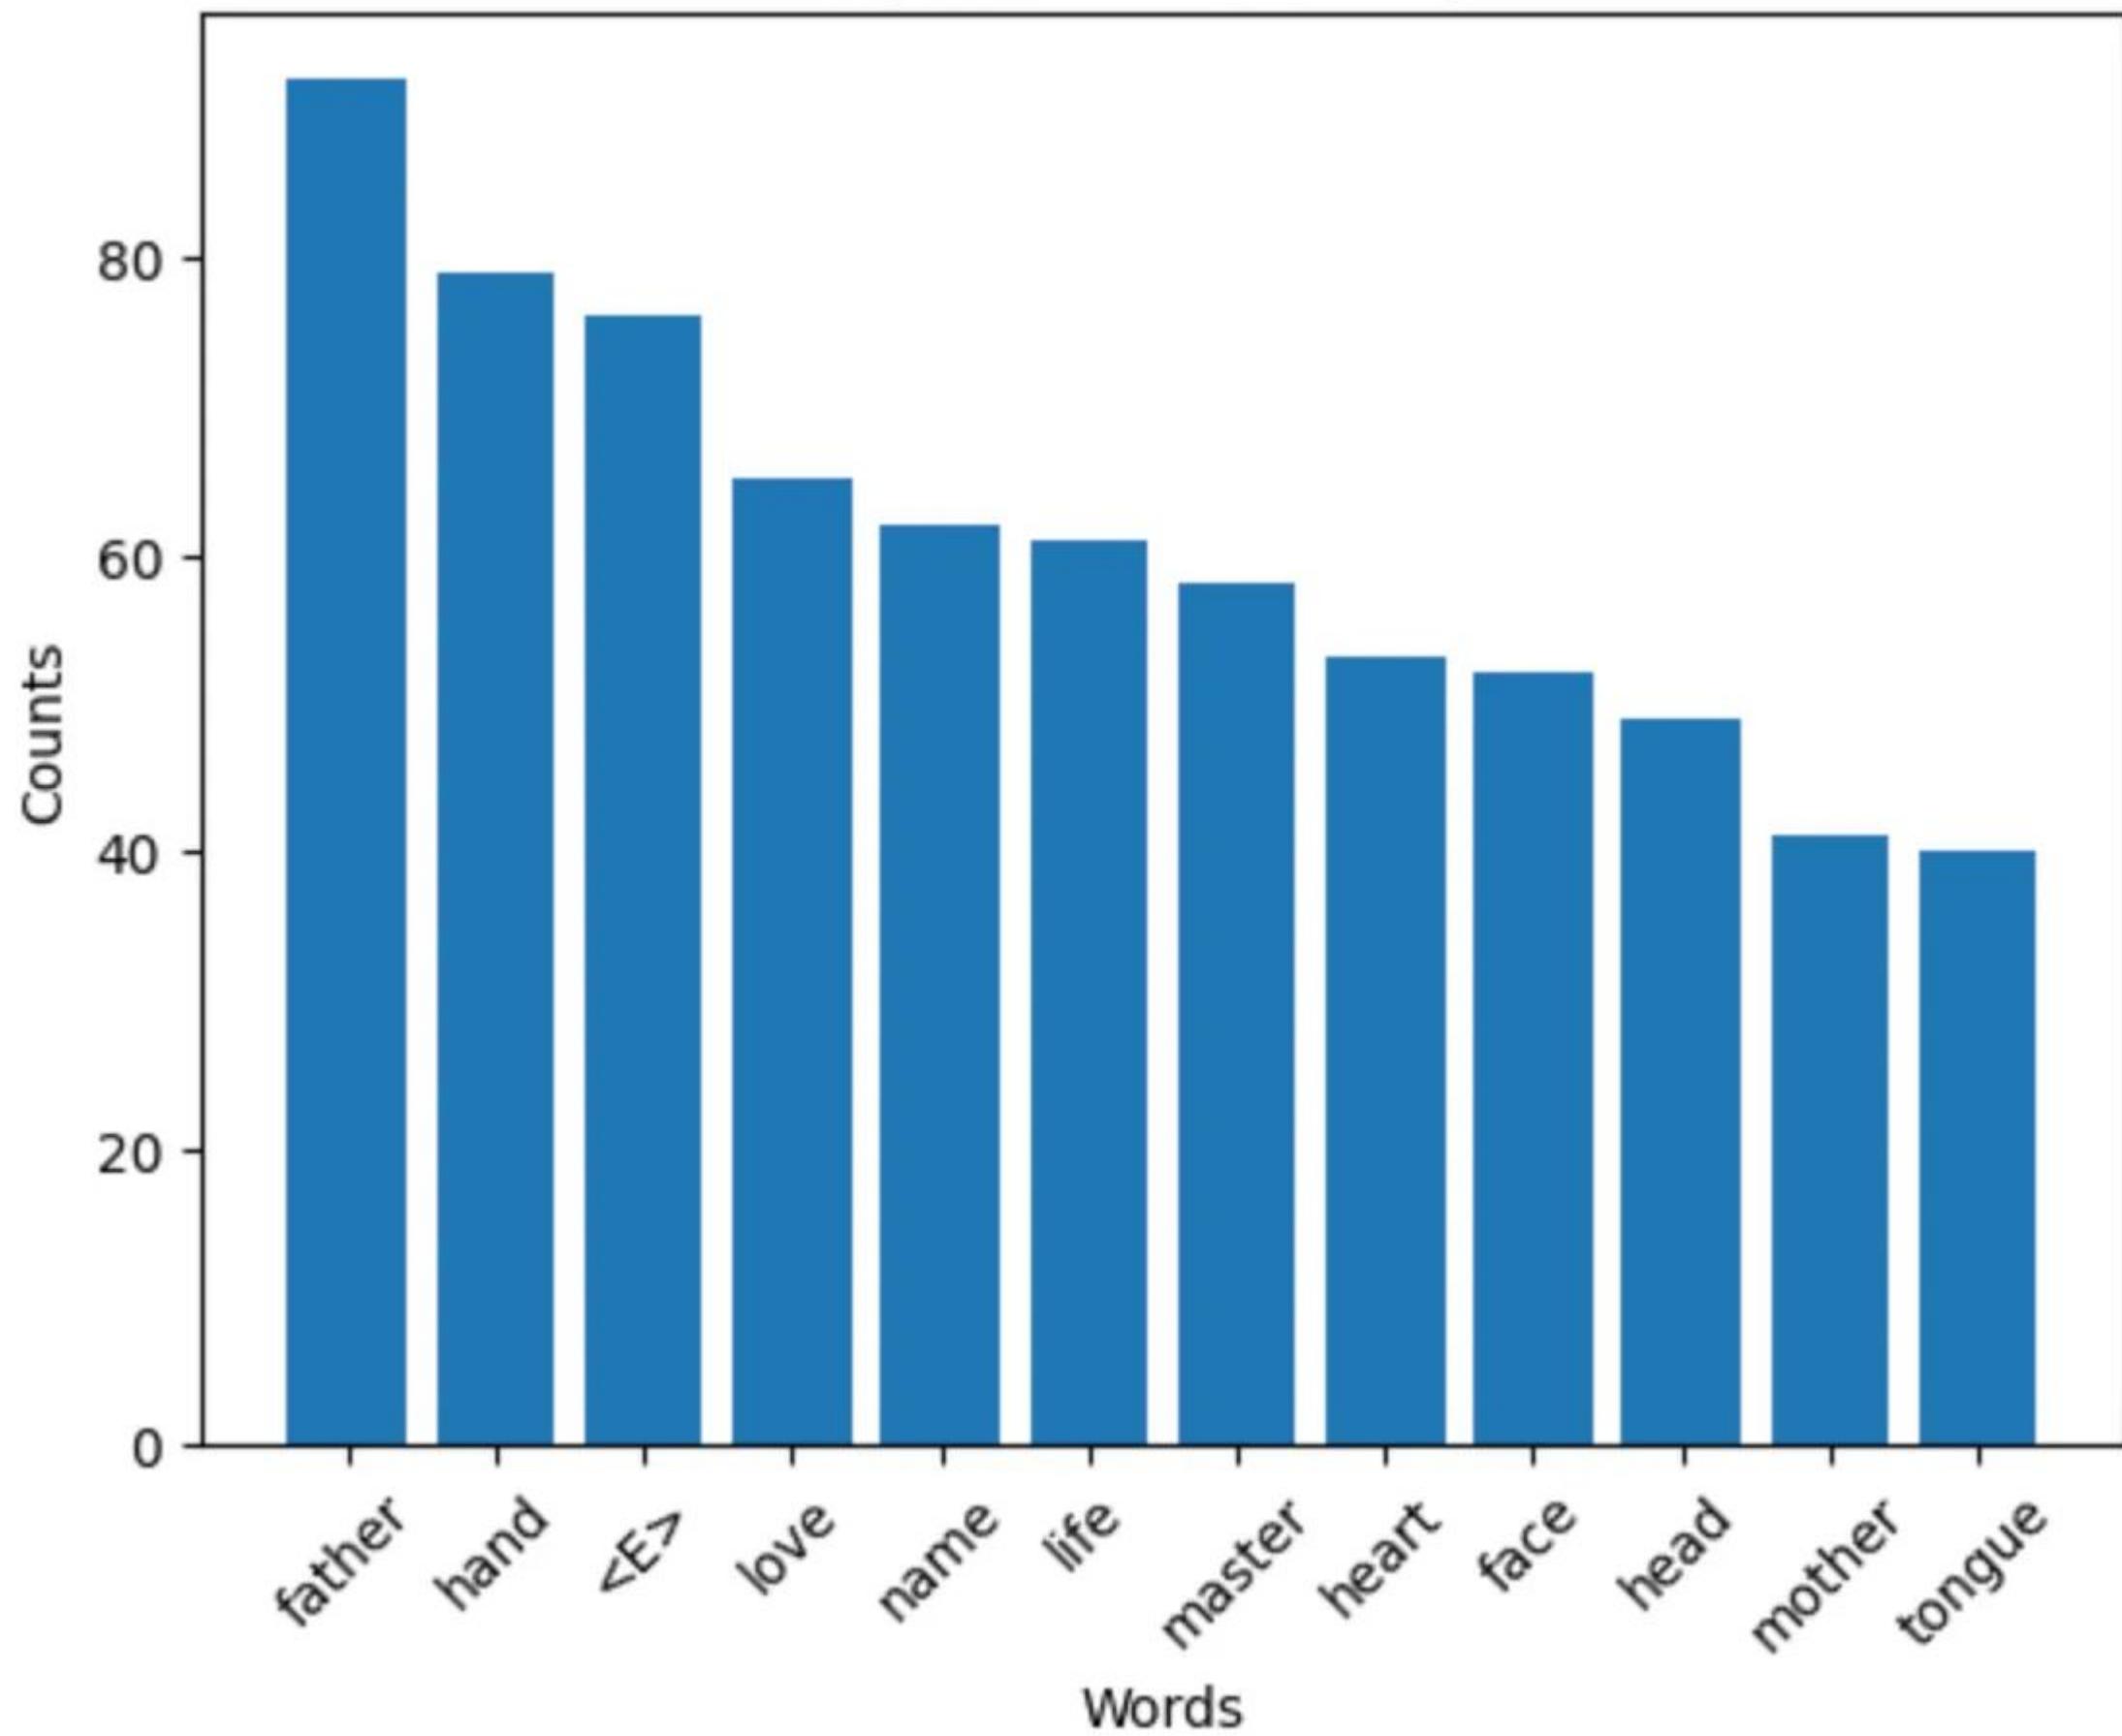
Let's say we want to start writing
like Shakespeare!

Inference

Now our model is trained.

Let's say we want to start writing
like Shakespeare!

> Thy

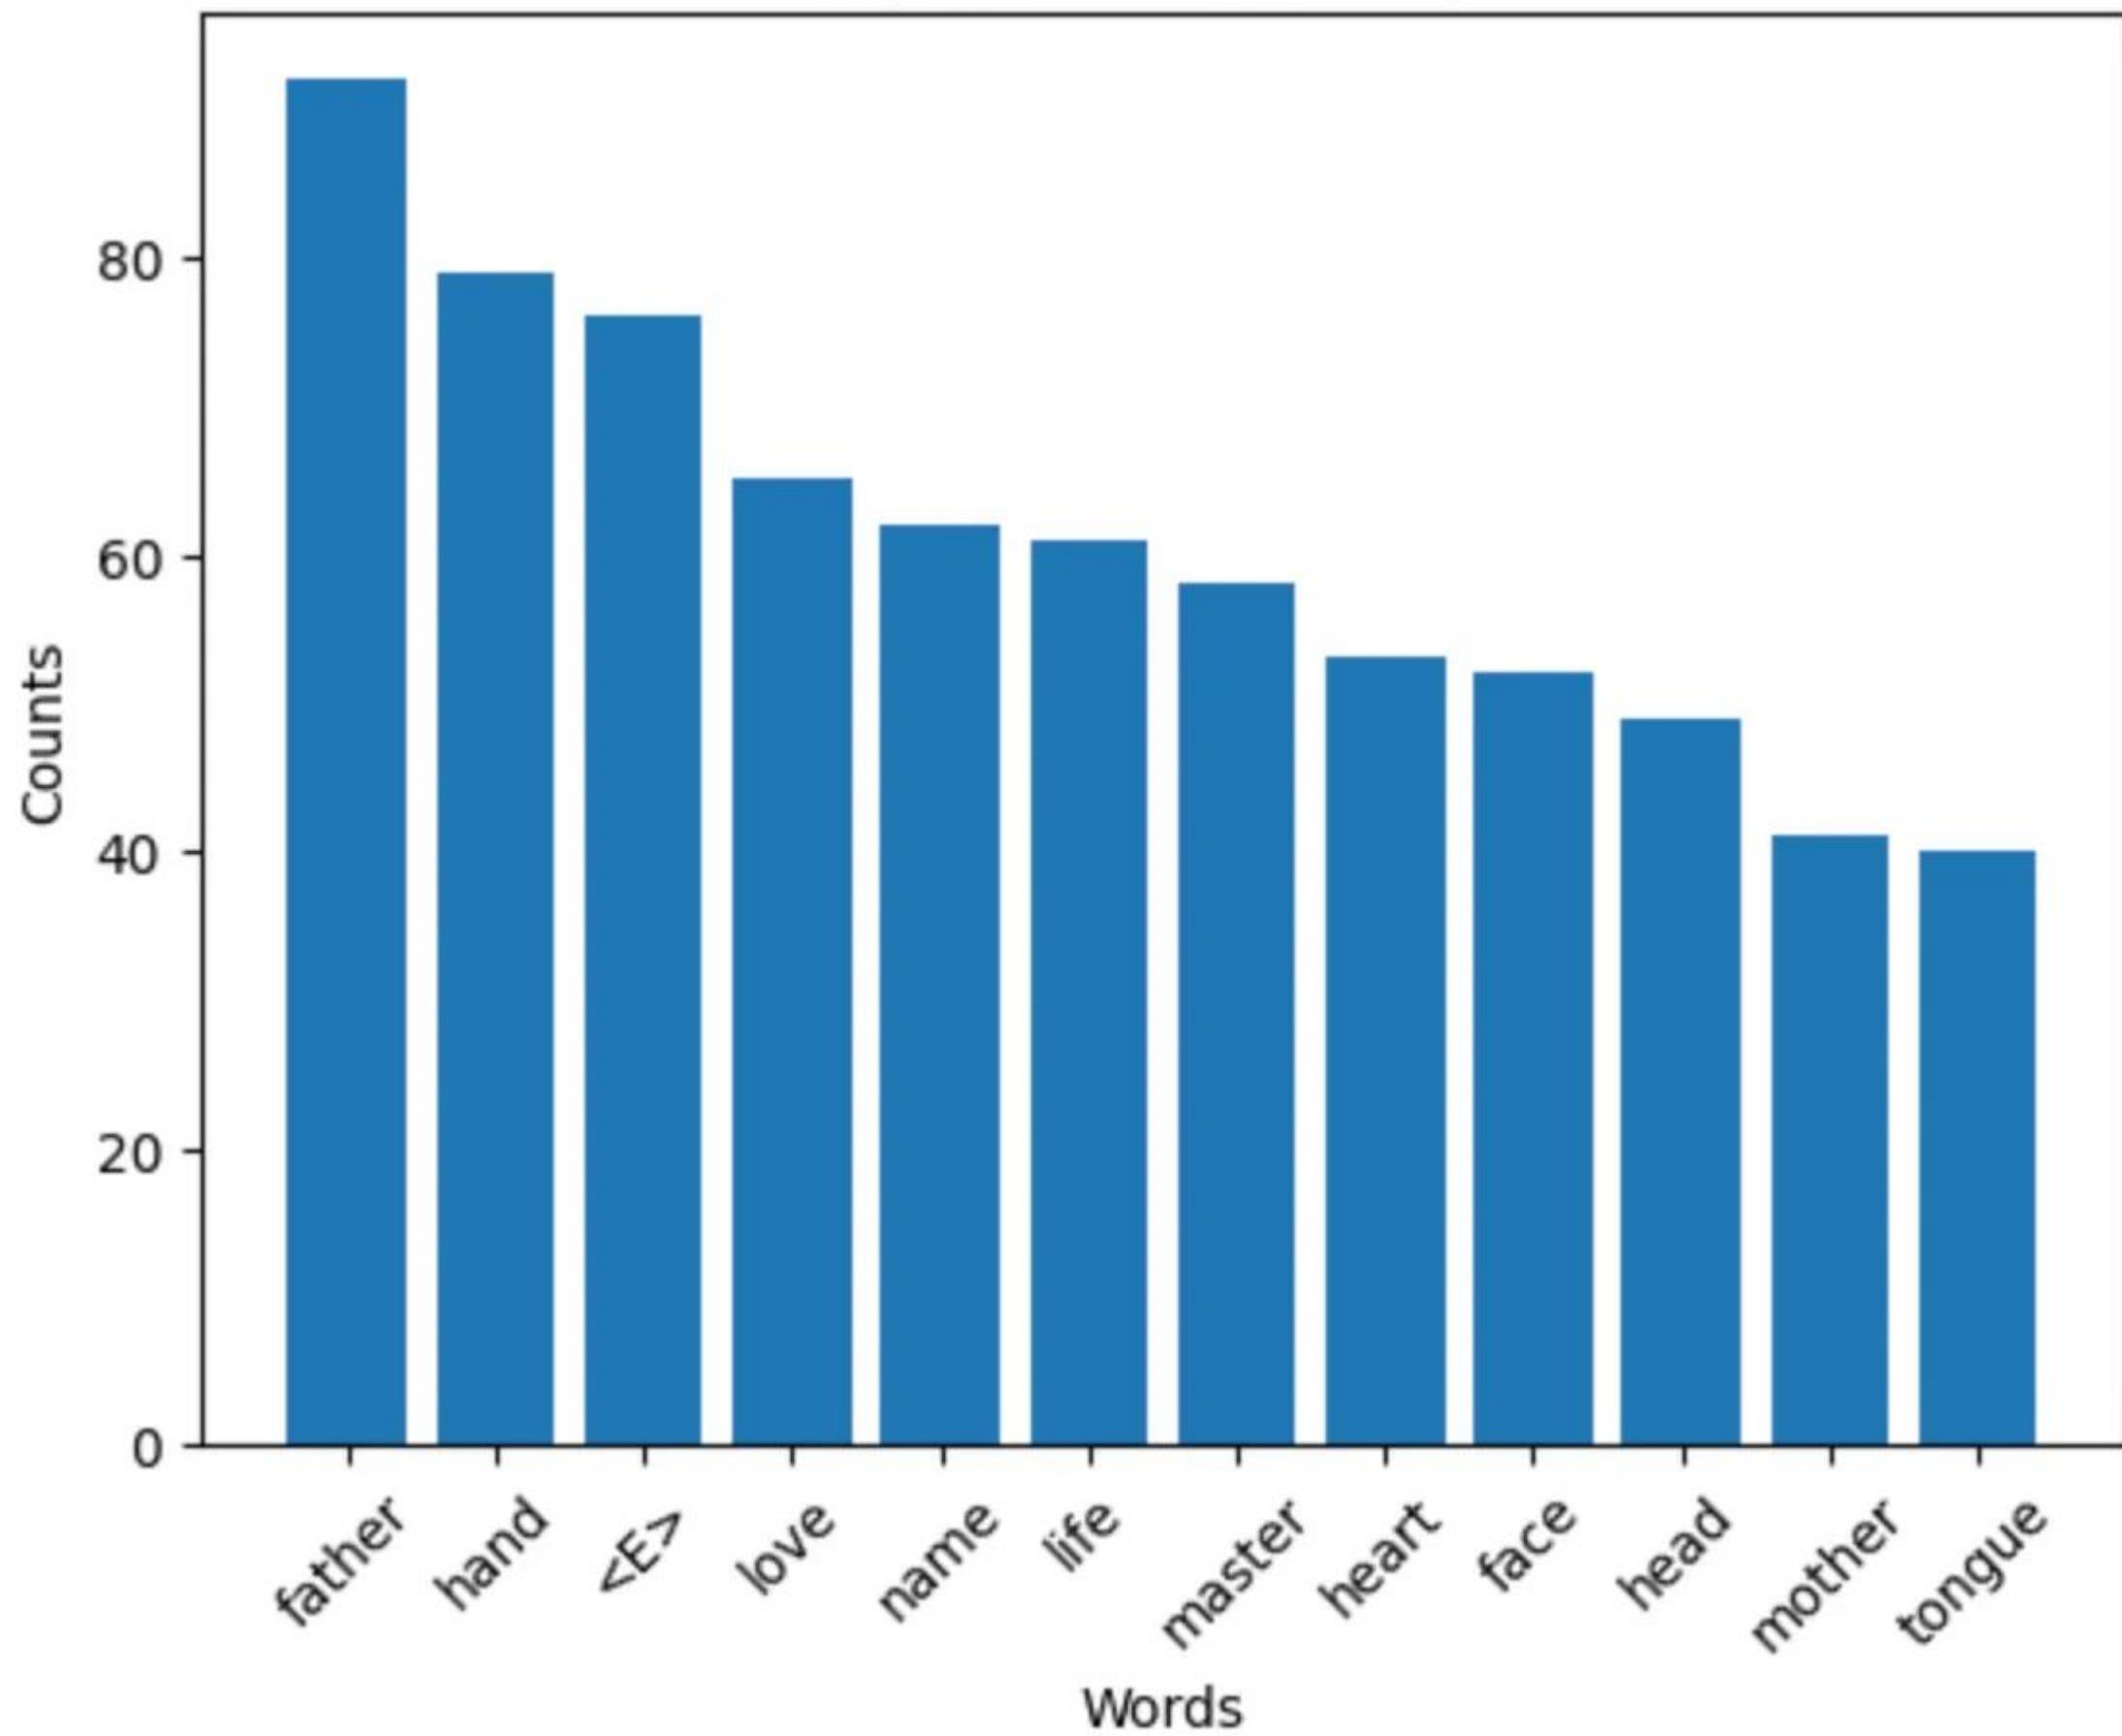


Inference

Now our model is trained.

Let's say we want to start writing
like Shakespeare!

> Thy father



Inference

Now our model is trained.

Let's say we want to start writing
like Shakespeare!

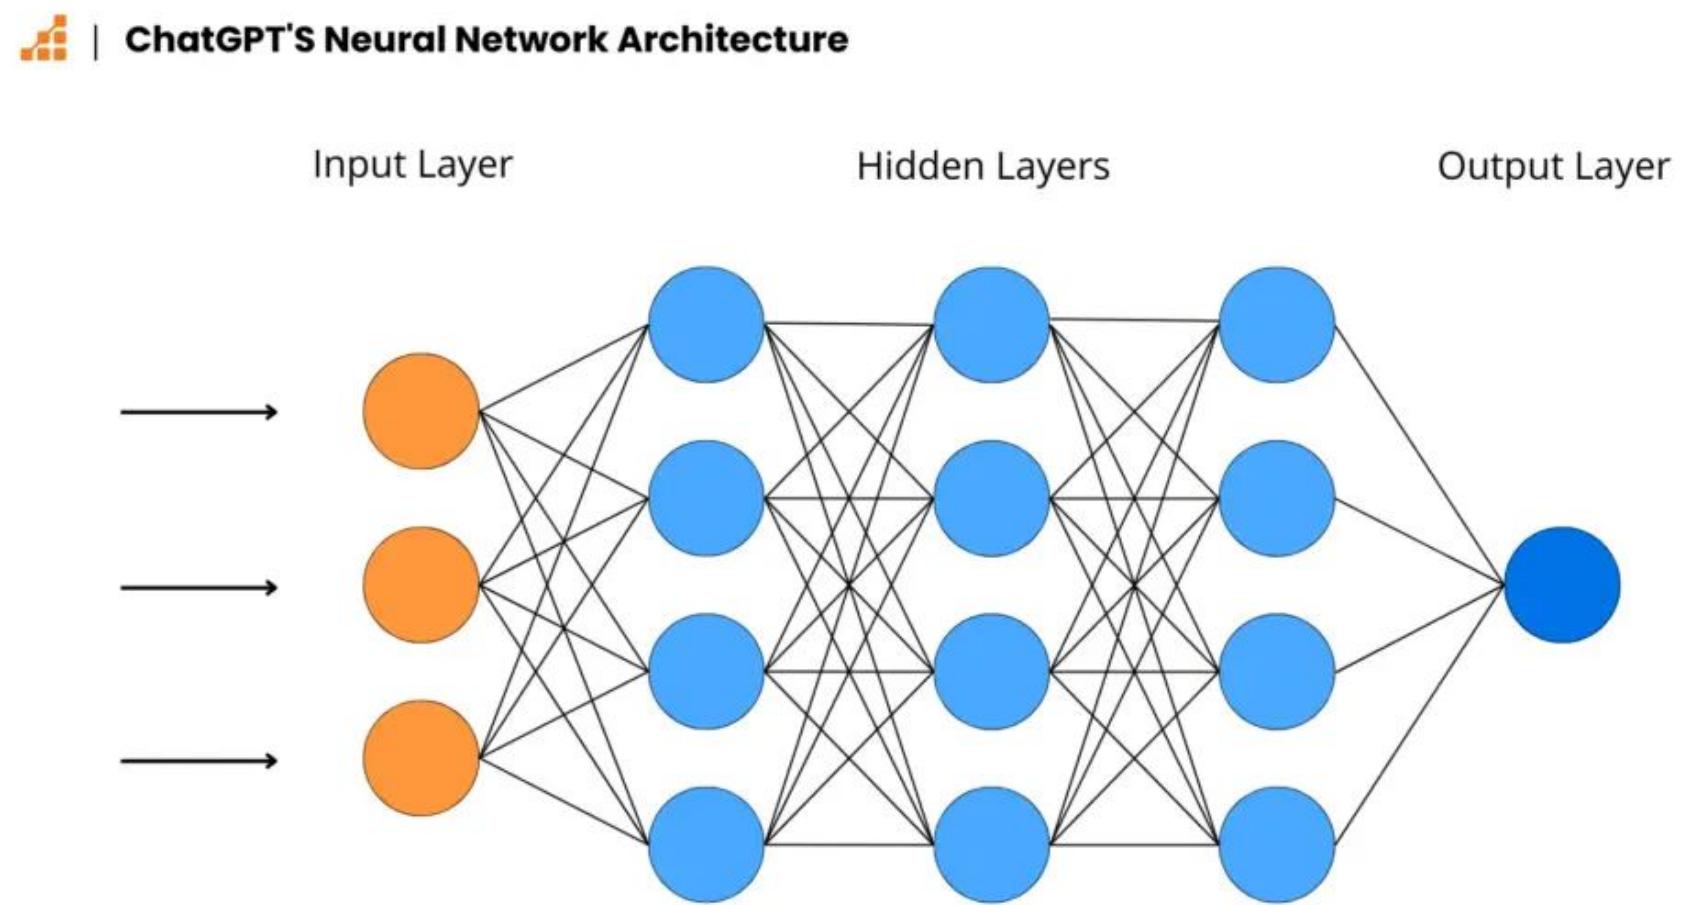
> Thy father and there were...

Does it work?

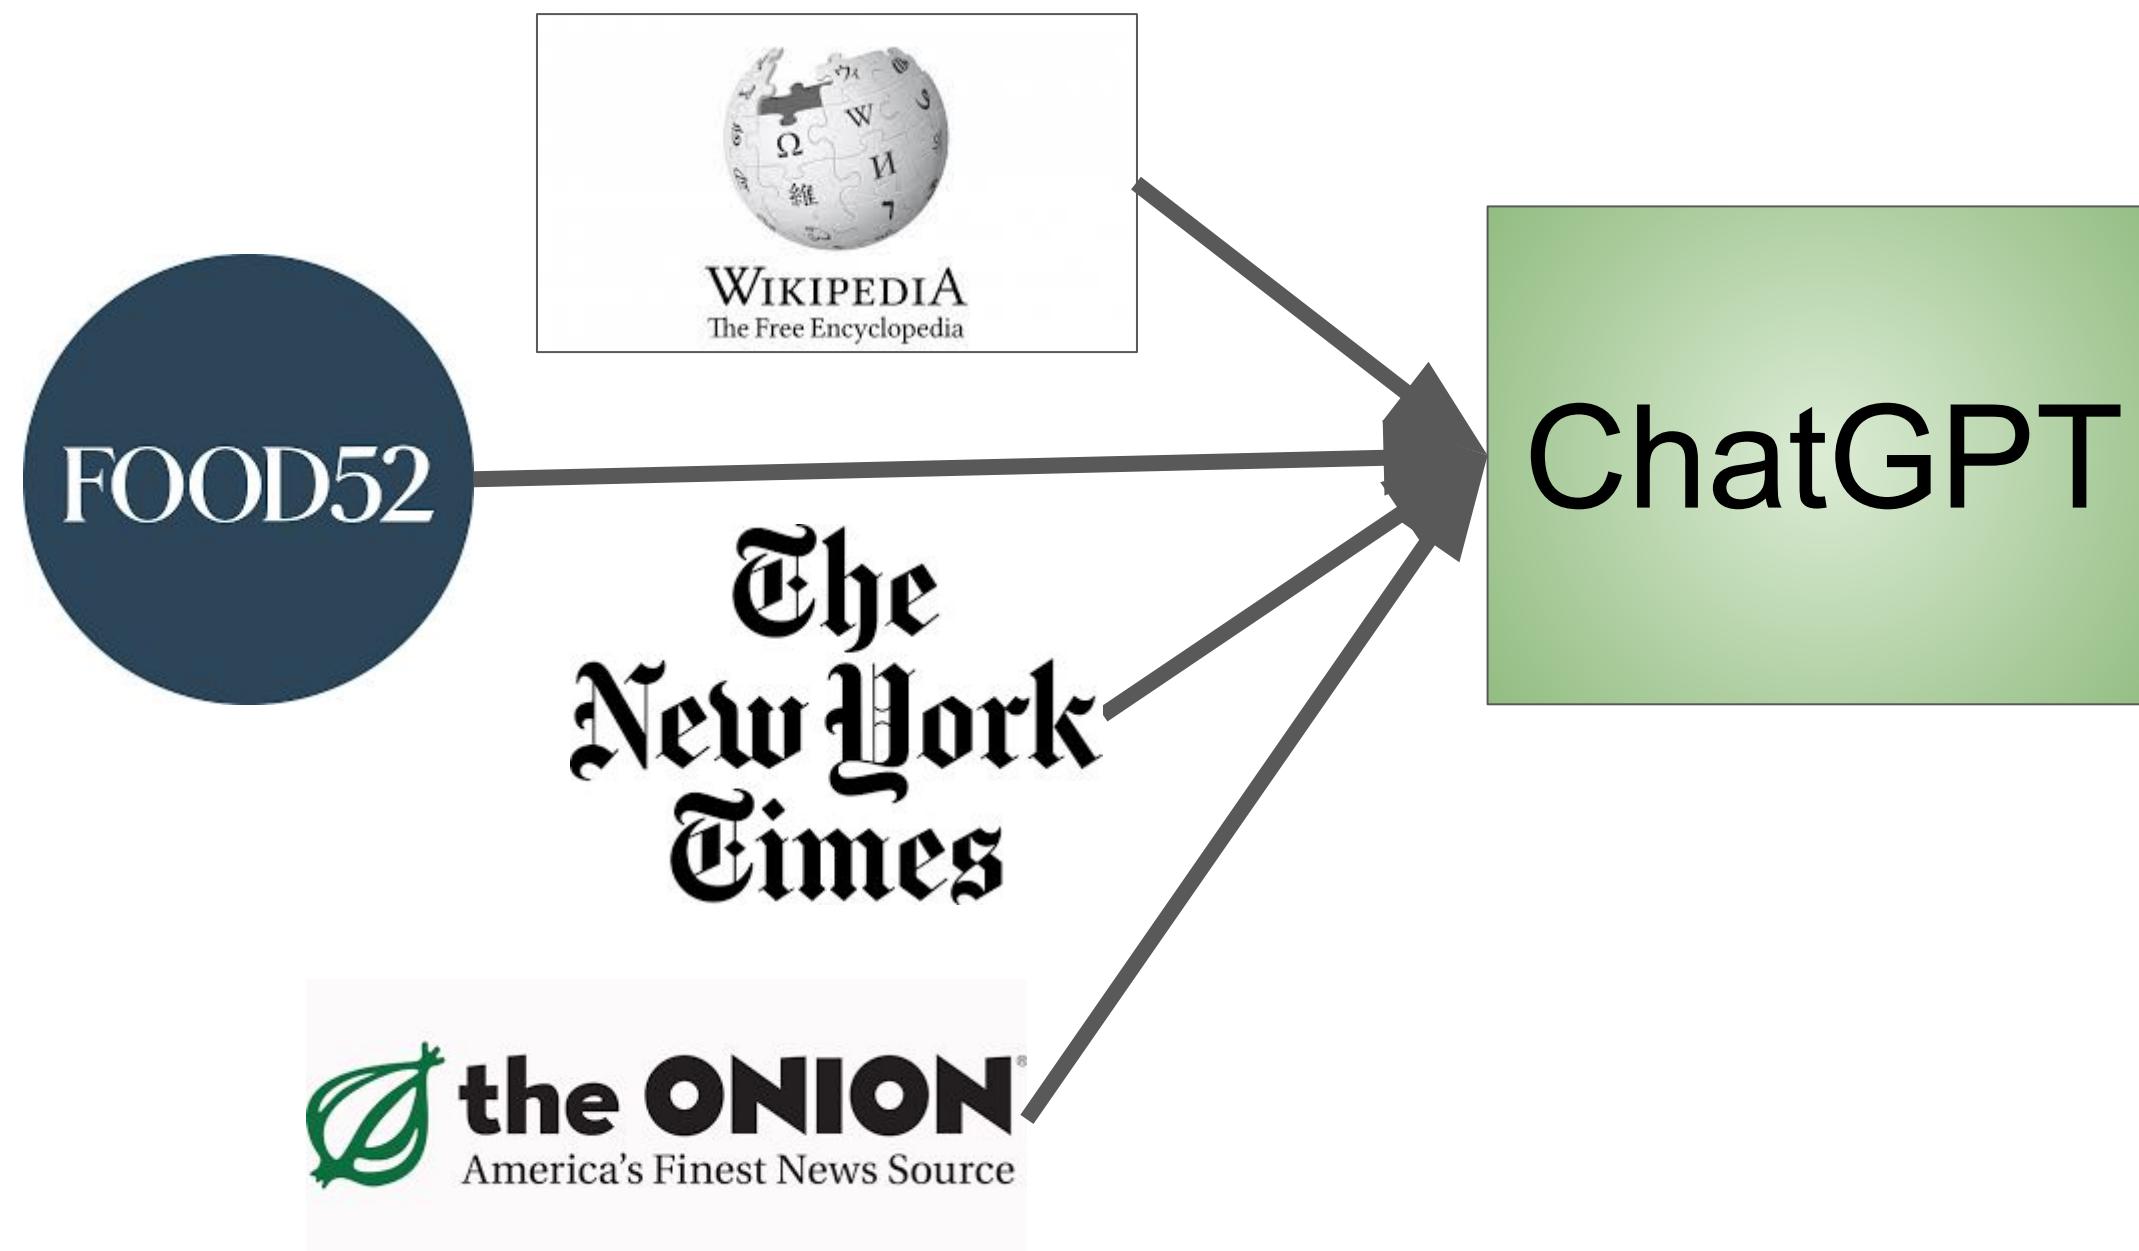
Bigrams are very simplistic.

They lack *context*.

So, actual language models do more than bigrams. ChatGPT uses a neural network.



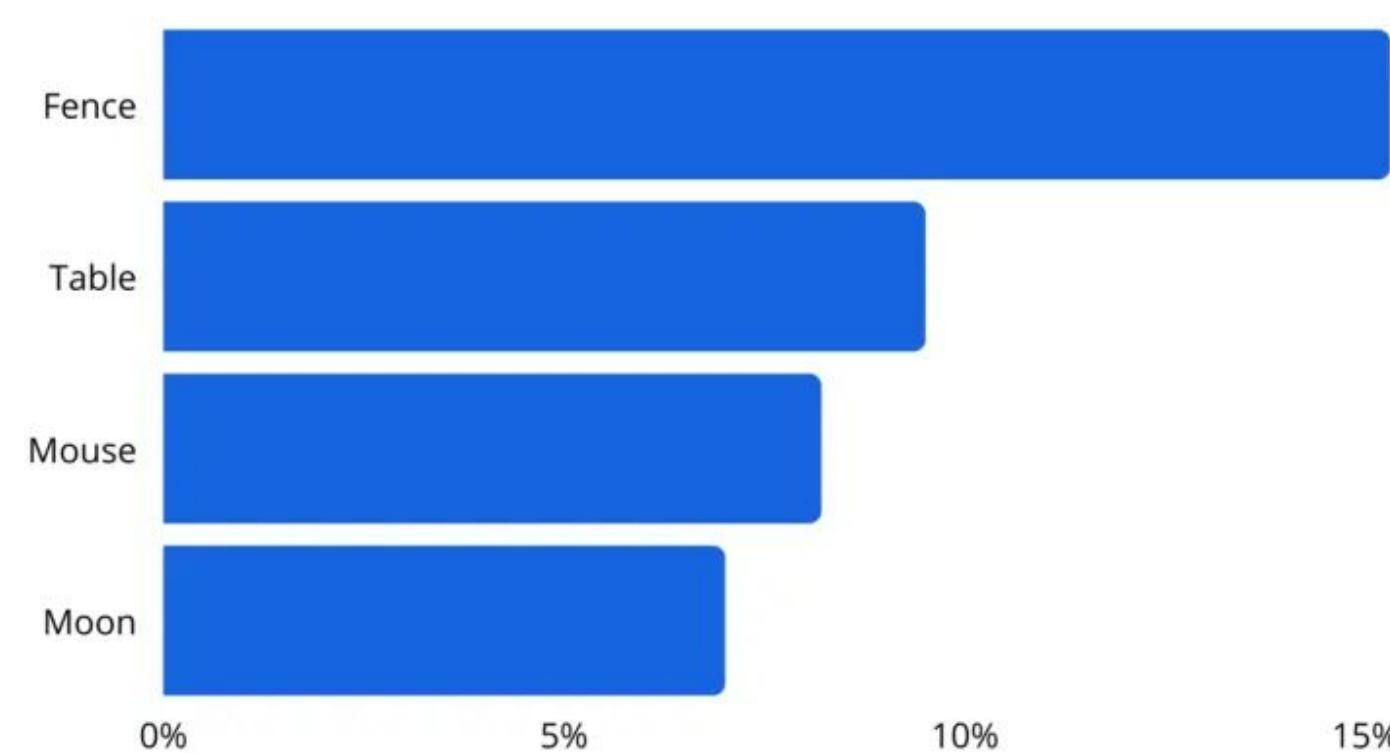
What goes into the Input Layer?



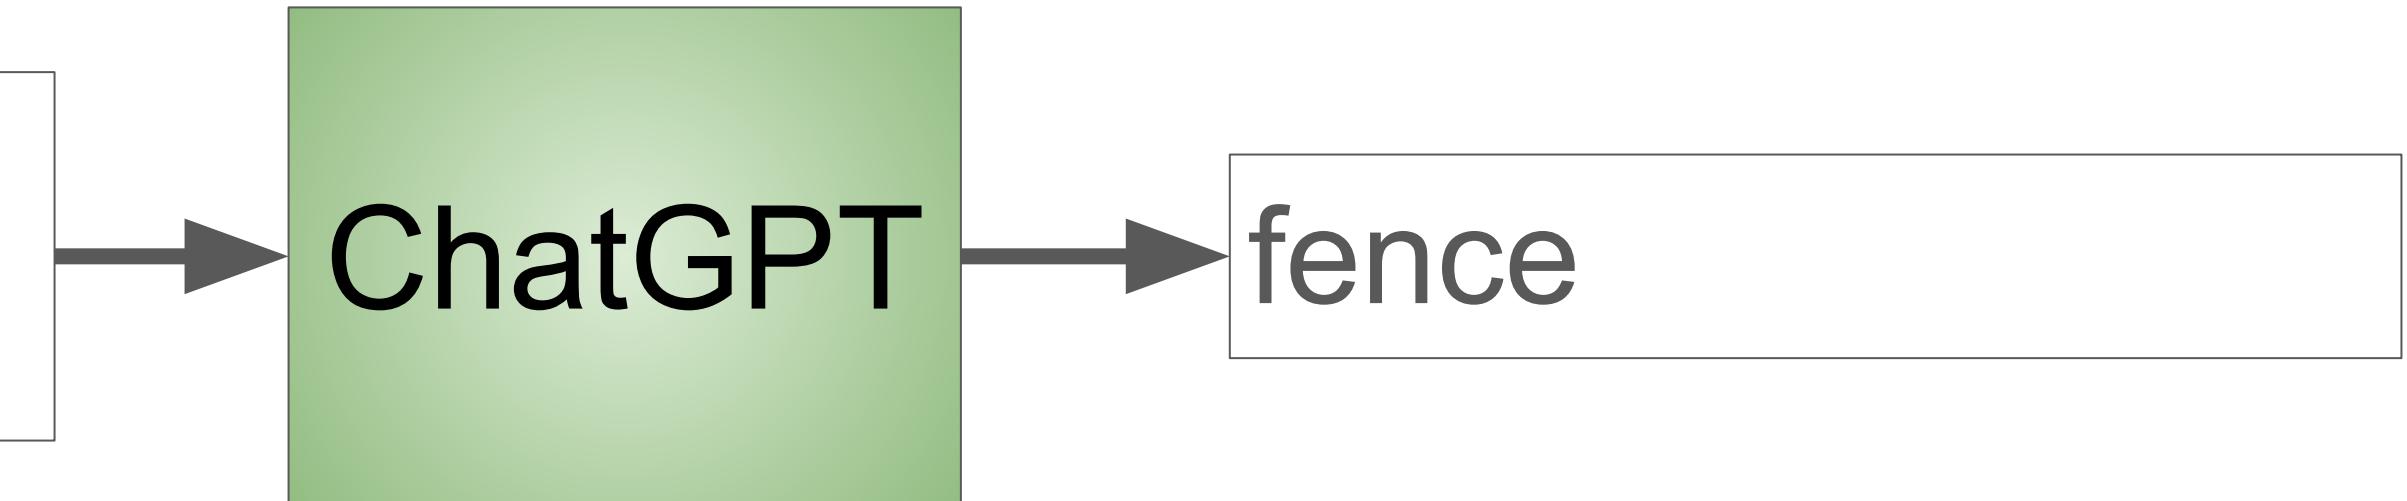
The cat jumped
over the...

ChatGPT

Prompt: "The cat jumped over the..."



The cat jumped
over the...



Some Neural Networks are Probabilistic

 Complete the sentence, "the cat jumped over the..."

 ...fence." thumb up thumb down

 Complete the sentence, "the cat jumped over the..."

 ...lazy dog." thumb up thumb down

 Complete the sentence, "the cat jumped over the..." edit

 ...candlestick." thumb up thumb down

Some Neural Networks are Probabilistic

Complete the sentence, "the cat jumped over the..."

...fence.  

Complete the sentence, "the cat jumped over the..."

...lazy dog.  

Complete the sentence, "the cat jumped over the..." 

...candlestick.  

Some Neural Networks are Probabilistic

Complete the sentence, "the cat jumped over the..."

...fence.

Complete the sentence, "the cat jumped over the..."

...lazy dog.

Complete the sentence, "the cat jumped over the..."

...candlestick.

Key Takeaway 1: Garbage in, Garbage out

Lots of hidden complexities and risks

Bias

Hallucinations (a.k.a. Correctness)

Privacy

Bias

Sources of bias

Training data

Algorithmic bias (e.g., color is more important than shape in classifying fruit)

Cognitive bias (e.g., based on societal factors)

Bias is common

Insight - Amazon scraps recruiting tool that showed women

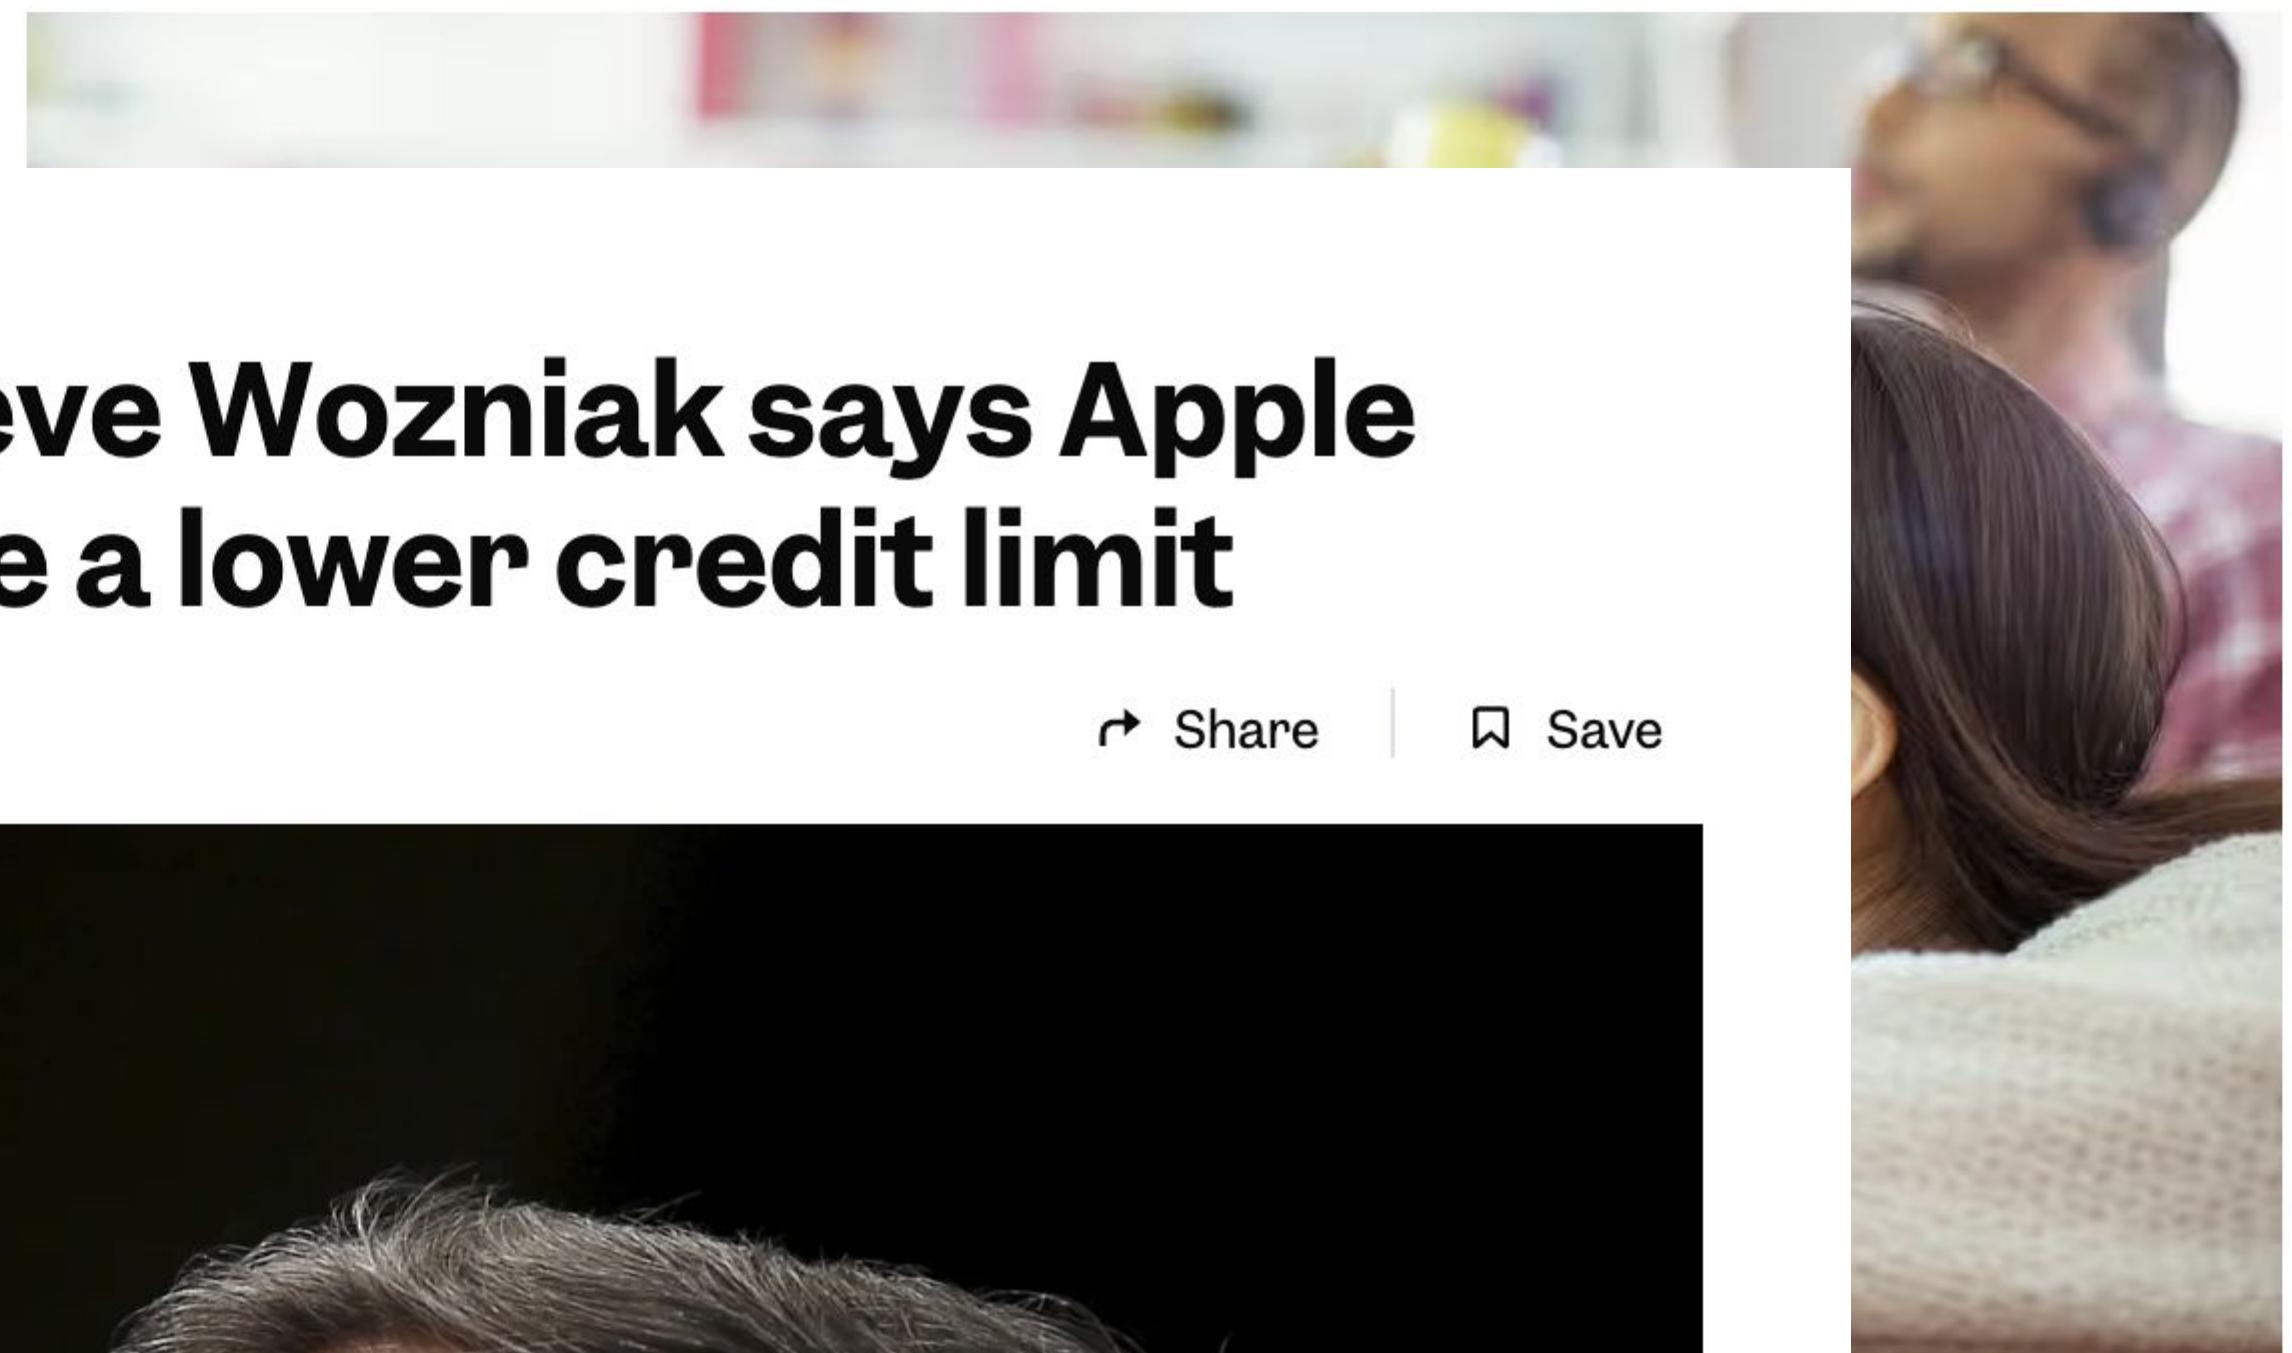
TECH

Apple cofounder Steve Wozniak says Apple Card offered his wife a lower credit limit

Isobel Asher Hamilton Nov 11, 2019, 7:21 AM ET

↗ Share

↪ Save



Women less likely to be shown ads for high-paid jobs on Google, study shows

Automated testing and analysis of company's advertising system reveals male job seekers are shown far more adverts for high-paying executive jobs

How can we detect bias or prevent bias?

Explicitly test for bias using datasets designed for that purpose

Train models on unbiased datasets

Reinforcement learning (use rewards and punishments in training)

Hallucinations

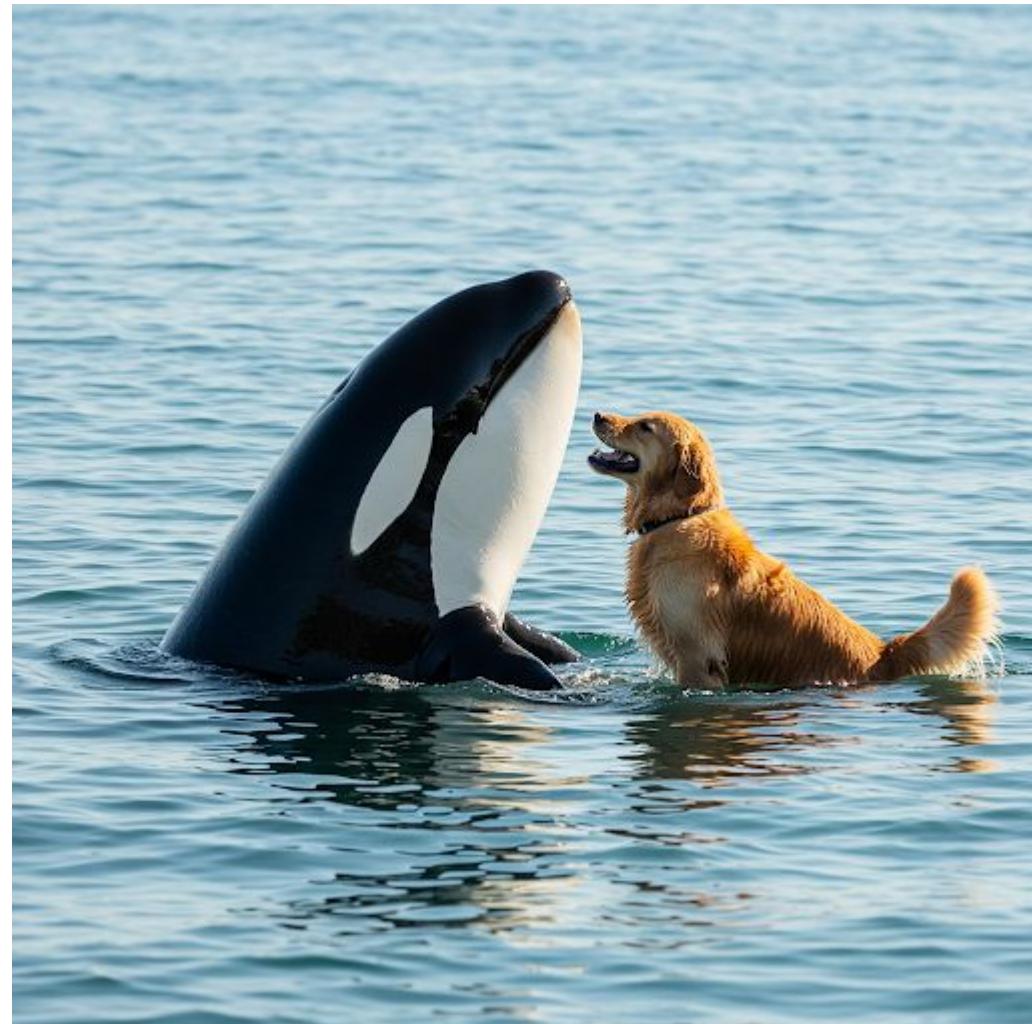
Hallucinations

It's when an AI creates something that isn't real.

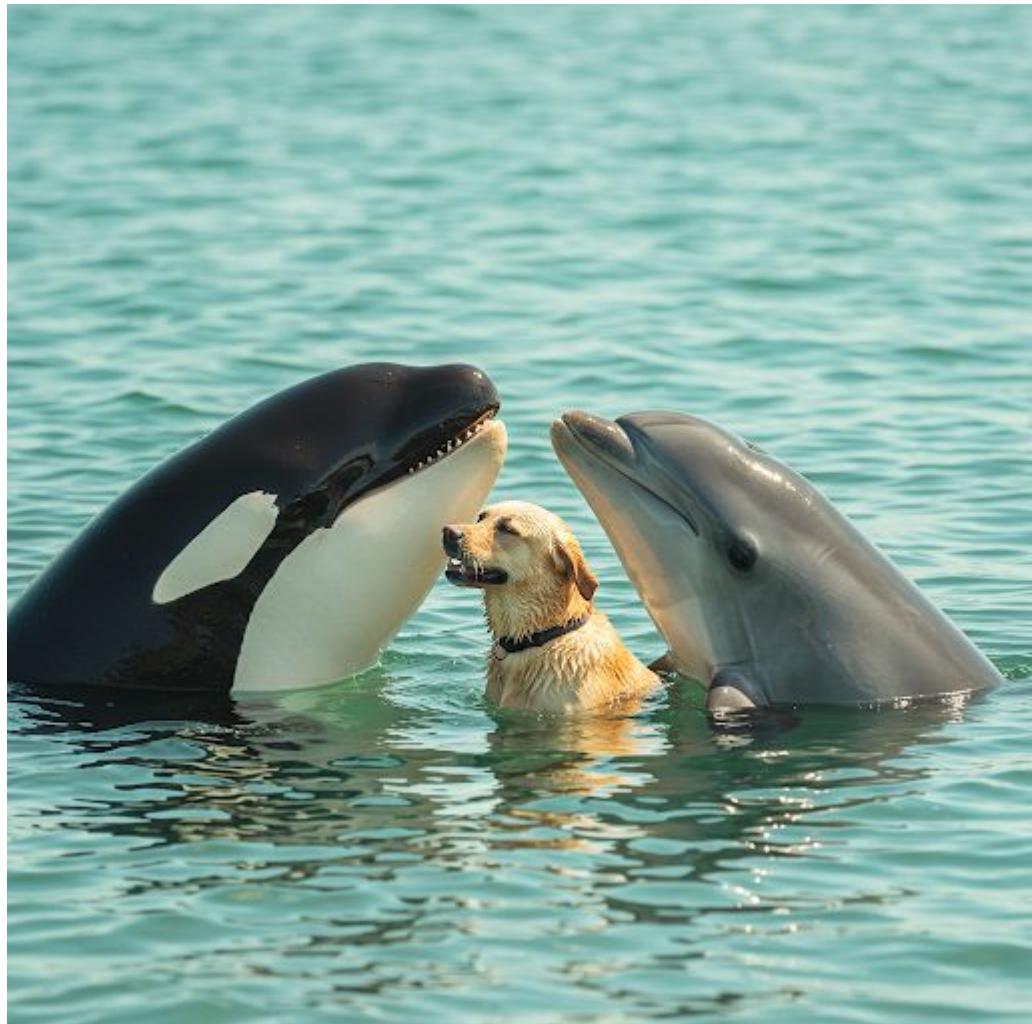


Tell me everything you know about orcas.





AI



AI



Real

Hallucinations

It's when an AI creates something that isn't real.

The screenshot shows a conversational AI interface. A user input box at the top contains the question "S what is the world record for crossing the english channel entirely on foot". Below it, a response from a green AI icon says: "The world record for crossing the English Channel entirely on foot is held by Christof Wandratsch of Germany, who completed the crossing in 14 hours and 51 minutes on August 14, 2020. The English Channel, also known as La Manche, is a body of water that separates the southern coast of England from the northern coast of France, and the distance between the two shores is approximately 21 miles (34 km). Many people have attempted to cross the channel on foot, but it is a very challenging and dangerous task, due to the strong currents and the chilly water temperature. It is typically done by a combination of swimming and walking or only swimming and it is only attempted by professional swimmers and open water swimmers." The first sentence of the response is highlighted with a red box, and the entire response area has a red border.

S what is the world record for crossing the english channel entirely on foot

 The world record for crossing the English Channel entirely on foot is held by Christof Wandratsch of Germany, who completed the crossing in 14 hours and 51 minutes on August 14, 2020. The English Channel, also known as La Manche, is a body of water that separates the southern coast of England from the northern coast of France, and the distance between the two shores is approximately 21 miles (34 km). Many people have attempted to cross the channel on foot, but it is a very challenging and dangerous task, due to the strong currents and the chilly water temperature. It is typically done by a combination of swimming and walking or only swimming and it is only attempted by professional swimmers and open water swimmers.

Regenerate response

Hallucinations

It's when an AI creates something
that isn't real.

More insidiously...

No

did the groundhog see his shadow this year?



did the groundhog see his shadow this year?



ALL COPILOT NEWS IMAGES VIDEOS MAPS

Did not see his shadow

According to 3 sources

It's official: spring is right around the corner. Punxsutawney Phil did not **see his shadow** in front of thousands of spectators Friday morning at Gobbler's Knob in Punxsutawney, Pennsylvania, part of the annual **Groundhog Day** tradition celebrated in the U.S. and Canada.

[Did the groundhog see his shadow this year?](#)
usatoday.com

The famous groundhog Punxsutawney before **not seeing** his shadow was an early indicator of the coming spring.

[Groundhog Day 2025: Did Punxsutawney Phil...](#)
today.com

did the groundhog see his shadow this year 2025?



did the groundhog see his shadow this year 2025?



ALL COPILOT NEWS IMAGES VIDEOS MAPS

Punxsutawney Phil, the famed Pennsylvania groundhog, **saw his shadow on Sunday morning** [1](#) [2](#) [3](#).

He predicted six more weeks of winter for 2025 [1](#) [2](#) [3](#) [4](#).

Learn more:

[1 Groundhog Day 2025: Did Punxsutawney Phi...](#)

[2 Did Punxsutawney Phil see his shadow? Here's what we know.](#)

Key Takeaway 2: Ask multiple times, multiple ways

ChatGPT is not a calculator

$378 + 63 = 441$. Let's ask Gemini if it's correct...



ChatGPT is not a calculator

$378 + 63 = 441$. Let's ask Gemini if it's correct...



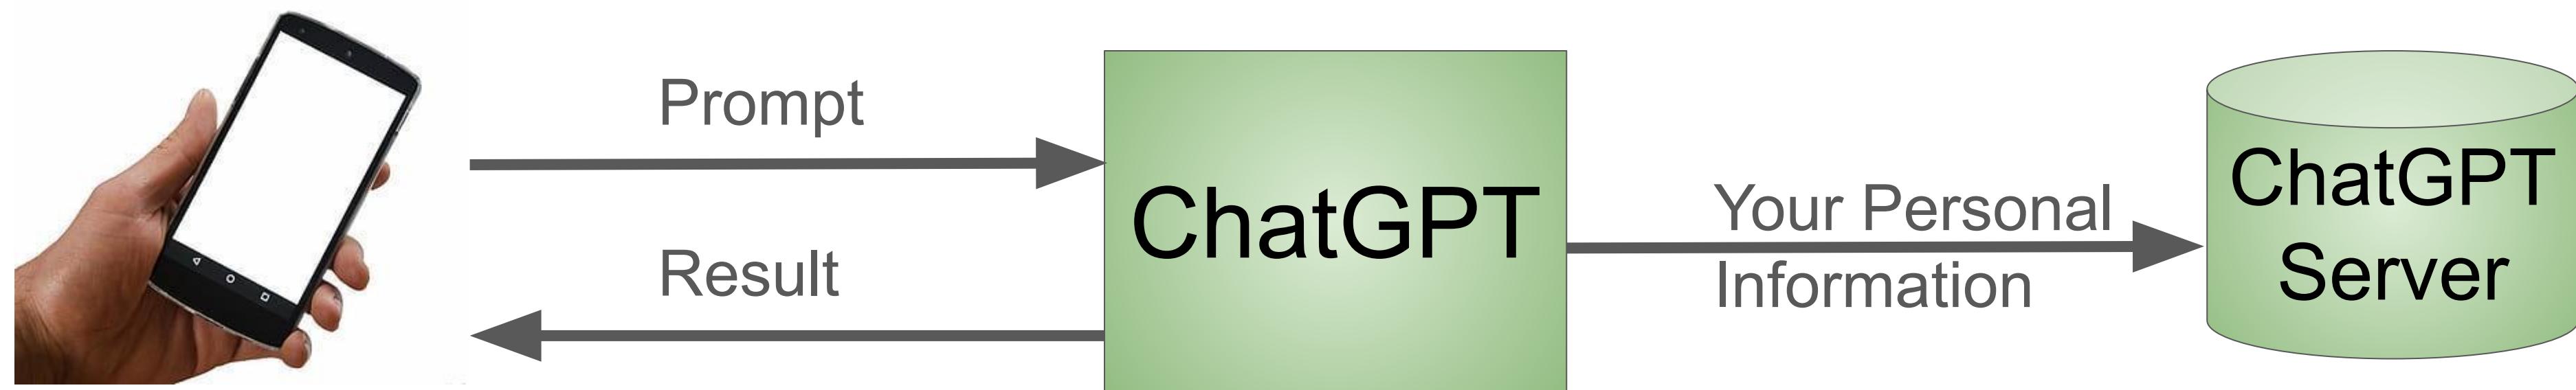
Oops, a typo! Or, my math was wrong?

Key Takeaway 3:
ChatGPT doesn't "think" or "compute"

Privacy

Privacy

Everything you tell an LLM is no longer private.



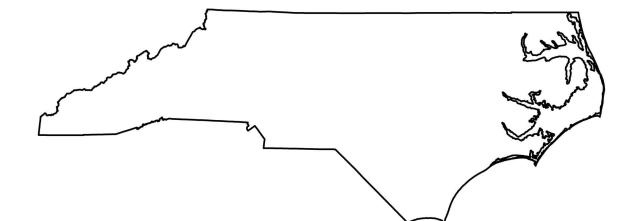
Summary

1. Garbage in, Garbage out: an AI is only as good as its data
2. Ask multiple times, multiple ways to detect potential hallucinations
3. AI doesn't "think" or "compute", it pattern-matches

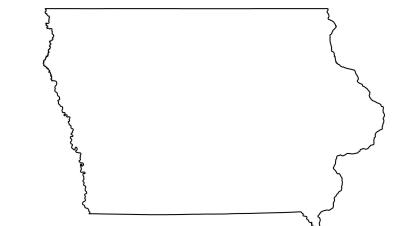
Enough about AI...
What about my research?

Who am I?

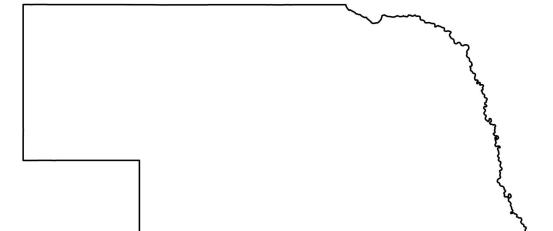
2016 – present: Associate Professor at NC State University



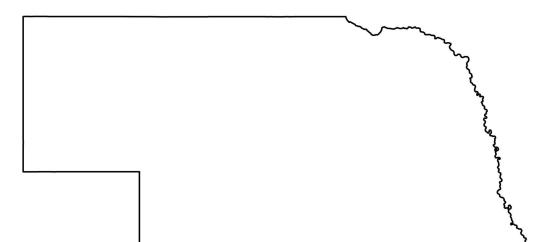
2013 – 2015: Assistant Professor at Iowa State University



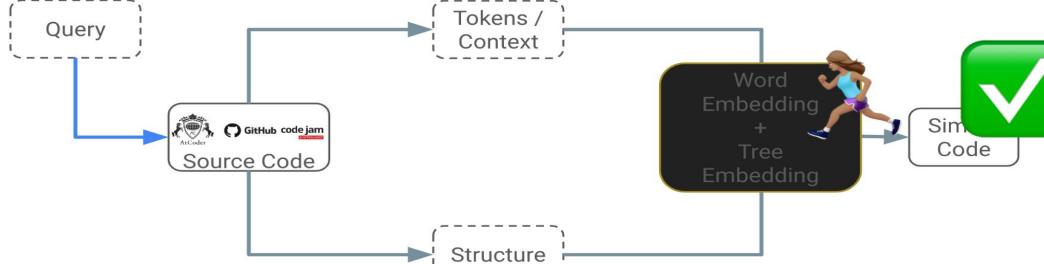
2008 – 2013: Graduate Student at University of Nebraska-Lincoln
M.S. in Computer Science
Ph.D. in Computer Science



2004 – 2008: Undergraduate Student at University of Nebraska-Lincoln
Major: Computer Science
Minor: Business Administration



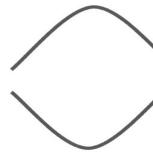
Code-to-code Search - In Practice



12

Comparative Comprehension

The cognitive activity of understanding how algorithms behave relative to each other



Participatory Design

Question

Additional context or details

Upvotes, Last updated, Reputation
Runnings

Upvotes, Last updated, Reputation
Runnings

Upvotes, Last updated, Reputation
Runnings

Run with [language] version x

Secondary responses

Code-first experience
Horizontal listing

16

Option A	Option B
How to build site	Descriptions from the authors
Search Options	Description...
Language, Version, Platform	...
OPTION A	OPTION B
GitHub A	GitHub B
...	...
public static boolean allSame(String[] x)	public static boolean allSame(String[] x)
// this initial capacity can be found	// this initial capacity can be found
int max = 0;	int max = 0;
for (int i = 0; i < x.length(); i++)	for (int i = 0; i < x.length(); i++)
if (x[i].length() > max)	if (x[i].length() > max)
max = i;	max = i;
return max;	return max;
OPTION Z	Encouraging interactive examples



41



FIND IT



THINK IT



CHOOSE IT

Questions?

Credits

<https://newsletter.pnote.eu/p/shakespeare-language-model>

<https://www.scalablepath.com/machine-learning/chatgpt-architecture-explained>